Introduction to R Programming
Data Frames

Pedro Fonseca

07 Maio 2020

The c() Operator

▶ Vectors are essential building blocks for handling multiple items in R.

▶ To create vectors use the *combine* operator:

```
x <-  -1

y <-  10

z <-  c(1, 6, 3)

(myvec <- c(1, 3, 1, 42))
```

```
## [1]  1  3  1 42
```

```
(myvec2 <- c(myvec, x, y ,z))
```

```
## [1]  1  3  1 42 -1 10  1  6  3
```

```
(myvec3 <-  c(myvec, 1, 2))
```

Subsetting

Get the first element:

```
myvec[1]
```

```
## [1] 1
```

Get the second element:

```
myvec[2]
```

```
## [1] 3
```

Subsetting

Get the first three elements:

```
myvec[1:3]
```

```
## [1] 1 3 1
```

Omit the first element:

```
myvec[-1]
```

```
## [1]  3  1 42
```

Omit more than one element:

```
myvec[-c(1,2)]
```

```
## [1]  1 42
```

Overwriting

Substitute an element:

```
myvec[3] <- 6
myvec
```

```
## [1]  1  3  6 42
```

Substitute more than one element:

```
myvec[c(2,3,4)] <- c(2,3,4)
```

Functions to Generate Vectors

```r
1:10
```

```
## [1]  1  2  3  4  5  6  7  8  9 10
```

```r
5:1
```

```
## [1] 5 4 3 2 1
```

```r
seq(1, 10)
```

```
## [1]  1  2  3  4  5  6  7  8  9 10
```

```r
seq(from = 18, to = 27, by = 3)
```

```
## [1] 18 21 24 27
```

Functions to Generate Vectors

```r
rep(x = 1, times = 4)
```

```
## [1] 1 1 1 1
```

```r
rep(c(3, 6), times = 3)
```

```
## [1] 3 6 3 6 3 6
```

```r
rep(c(3, 62, 8.3), each = 2)
```

```
## [1]  3.0  3.0 62.0 62.0  8.3  8.3
```

```r
rep(c(3, 6), times = 3, each = 2)
```

```
##  [1] 3 3 6 6 3 3 6 6 3 3 6 6
```

Sorting the Elements of a Vector

Sorting a vector in increasing or decreasing order:

```
myvec2 <- c(1, 3, 1, 42, -5, 10, -50)
myvec2
```

```
## [1]   1   3   1  42  -5  10 -50
```

```
sort(myvec2)
```

```
## [1] -50  -5   1   1   3  10  42
```

```
sort(myvec2, decreasing = TRUE)
```

```
## [1]  42  10   3   1   1  -5 -50
```

Sorting the Elements of a Vector

```r
sort(c(2.5, -1, -10, 3.44))
```

```
## [1] -10.00  -1.00   2.50   3.44
```

```r
sort(c(2.5,-1,-10,3.44), decreasing = TRUE)
```

```
## [1]   3.44   2.50  -1.00 -10.00
```

```r
sort(c(2.5,-1,-10,3.44), TRUE)
```

```
## [1]   3.44   2.50  -1.00 -10.00
```

Some Statistical Functions

rnorm(n) generates n pseudo-random numbers from a normal distribution (default: $\mu = 0$, $\sigma = 1$)

```
rnorm(3)
```

```
## [1] -1.2067290 -0.6803377 -0.2182001
```

```
rnorm(4, mean = 5, sd = 2)
```

```
## [1] 4.841469 5.639049 7.480188 6.494767
```

Other functions related do the normal distribution:

▶ dnorm (density), pnorm (distribution function), qnorm (quantile function).

▶ Equivalent functions are available for the most commonly used probability distributions: F, t-student, Uniform, Poisson. . .

The set.seed Function

- Functions like rnorm, rpois and runif generate pseudo-random numbers.

- This means that you and I will get different results when using these functions.

- Solution: use the set.seed function.

The set.seed Function

Try this command many times:

```
rnorm(2)
```

[1] 1.0358944 -0.7403707

Each time you will get a different output. Now try this:

```
set.seed(123)
rnorm(2)
```

[1] -0.5604756 -0.2301775

You will get the same output every time. The argument of *set.seed* is irrelevant as long as we all use the same value.

Vectorized Operations

One of the main advantages of R is vectorized calculation. This means that:

- ▶ Most R functions accept vectors as inputs;
- ▶ Vector arithmetic is performed element-wise by default.

Vectorized Operations

- Vectorization calculation is a huge advantage efficiency and parsimony.

- Vectorization also makes code easier to write and read.

Vectorized Operations

```
x <- c(1, 2, 3)
y <- c(0.5, 0.5, 0.5)
1/x
```

```
## [1] 1.0000000 0.5000000 0.3333333
```

```
3+y
```

```
## [1] 3.5 3.5 3.5
```

```
x+y
```

```
## [1] 1.5 2.5 3.5
```

Vectorized Operations

```r
x <- c(1, 2, 3)
y <- c(0.5, 0.5, 0.5)

x^y
```

```
## [1] 1.000000 1.414214 1.732051
```

```r
sqrt(x)
```

```
## [1] 1.000000 1.414214 1.732051
```

```r
1/1:3
```

```
## [1] 1.0000000 0.5000000 0.3333333
```

```r
seq(from = 2, to = 6, by = 2)/2
```

```
## [1] 1 2 3
```

Vectorized Operations

```r
x1 <- c(1, 5, 7)
x2 <- rep(1, times = 3)

log(x1)
```

```
## [1] 0.000000 1.609438 1.945910
```

```r
log(x1) - x2
```

```
## [1] -1.0000000  0.6094379  0.9459101
```

```r
x <- x1 + x2
x
```

```
## [1] 2 6 8
```

Rounding

round() rounds the values in its first argument to the specified number of decimal places (default 0):

```
set.seed(123)

z <- rnorm(3)
z
```

```
## [1] -0.5604756 -0.2301775  1.5587083
```

```
round(z, digits = 3)
```

```
## [1] -0.560 -0.230  1.559
```

```
round(z)
```

```
## [1] -1  0  2
```

Rounding

```
y <- c(3.271109, 3.374961, 2.313307, 4.837787)
round(y, 2)

## [1] 3.27 3.37 2.31 4.84
```

Statistical Functions

```
z
```

```
## [1] -0.5604756 -0.2301775  1.5587083
```

```
abs(z) # Absolut value
```

```
## [1] 0.5604756 0.2301775 1.5587083
```

```
max(z)
```

```
## [1] 1.558708
```

```
min(z)
```

```
## [1] -0.5604756
```

Statistical Functions

```
z
```

```
## [1] -0.5604756 -0.2301775  1.5587083
```

```
mean(z)
```

```
## [1] 0.2560184
```

```
median(z)
```

```
## [1] -0.2301775
```

```
sd(z)
```

```
## [1] 1.140186
```

Statistical Functions

```
z
```

```
## [1] -0.5604756 -0.2301775  1.5587083
```

```
var(z)
```

```
## [1] 1.300025
```

```
sum(z)
```

```
## [1] 0.7680552
```

```
quantile(z, 0.5)
```

```
##        50%
## -0.2301775
```

The which() Function

The which() function is useful to find which elements of a vector that verify a given condition:

```r
set.seed(123)
vec <- rnorm(n = 10, mean = 2, sd = 1)
round(vec, 2)
```

```
##  [1] 1.44 1.77 3.56 2.07 2.13 3.72 2.46 0.73 1.31 1.55
```

```r
(indexes <- which(vec > 2))
```

```
## [1] 3 4 5 6 7
```

```r
round(vec[indexes], 3)
```

```
## [1] 3.559 2.071 2.129 3.715 2.461
```

The which() Function

```
set.seed(123)

vec2 <- rpois(n = 10, lambda = 2)
which(vec2 == 2)
```

```
## [1]  3  7  9 10
```

```
(vec2 <- rpois(n = 10, lambda = 2))
```

```
##  [1] 5 2 3 2 0 4 1 0 1 5
```

```
which(vec2 == 2)
```

```
## [1] 2 4
```

The which() Function

```
set.seed(123)

vec2 <- rpois(n = 10, lambda = 2)
vec2
```

```
##  [1] 1 3 2 4 4 0 2 4 2 2
```

```
max(vec2)
```

```
## [1] 4
```

```
which(vec2 == max(vec2))
```

```
## [1] 4 5 8
```

The which() Function

The which function gives the positions of the elements of the vectors that verify the condition, not their values!

```
set.seed(123)

vec2 <- rpois(n = 10, lambda = 2)
vec2
```

```
## [1] 1 3 2 4 4 0 2 4 2 2
```

What are the actual values of *vec2* (not their positions) that verify the condition?

```
vec2[which(vec > 1)]
```

```
## [1] 1 3 2 4 4 0 2 2 2
```

The length() Function

```r
round(vec[which(vec > 2)], 3)
```

```
## [1] 3.559 2.071 2.129 3.715 2.461
```

Use length() to obtain the number of elements in a vector:

```r
length(vec)
```

```
## [1] 10
```

How many elements of vec are greater than 2?

```r
length(which(vec > 2))
```

```
## [1] 5
```

Trigonometric Functions

R trigonometric take radians as argument, not degrees:

- $\sin(\frac{\pi}{2})$:

```r
sin(pi/2)
```

```
## [1] 1
```

- $\cos(\pi)$:

```r
cos(pi)
```

```
## [1] -1
```

- $\tan(\frac{\pi}{3})$:

```r
tan(pi/3)
```

```
## [1] 1.732051
```

- cotangent$(\frac{\pi}{3})$:

```r
1/tan(pi/3)
```

Trigonometric Functions

Which value has a cosine = -1?

```
acos(-1)
```

[1] 3.141593

```
cos(pi)
```

[1] -1

Which value has a tangent = 0.5?

```
atan(0.5)
```

[1] 0.4636476

```
tan(0.4636476)
```

[1] 0.5

Trigonometric Functions

Trigonometric functions are also vectorized:

```r
(x <- seq(from = 0.25, to = 1, by = 0.25))
```

```
## [1] 0.25 0.50 0.75 1.00
```

```r
cos(x)
```

```
## [1] 0.9689124 0.8775826 0.7316889 0.5403023
```

```r
1/tan(x) # cotangent of x
```

```
## [1] 3.9163174 1.8304877 1.0734261 0.6420926
```

```r
cos(x)/sin(x) # cotangent of x
```

```
## [1] 3.9163174 1.8304877 1.0734261 0.6420926
```

Trigonometric Functions

R has many more trigonometric functions. Try:

```
?Trig
```

Recycling

What happens when we conduct calculations with two vectors of different length?

```
myvec <- c(1, 2, 3, 4)
myvec2 <- rep(0.5, times = 8)

myvec + myvec2
```

```
## [1] 1.5 2.5 3.5 4.5 1.5 2.5 3.5 4.5
```

Recycling

```r
myvec3 <- rep(0.5, times = 7)

myvec + myvec3

## Warning in myvec + myvec3: longer object length is not a
## object length

## [1] 1.5 2.5 3.5 4.5 1.5 2.5 3.5
```

Recycling

- When conducting operations that require input vectors to be of the same length, R automatically recycles, or repeats, the shorter one, until it is long enough to match the longer one.

- It will only throw an error message if the length of the shorter vector is not a multiple of the vector of the larger vector.

Named Vectors

We can also name the elements of a vector:

```
x <- c(x1 = 1, x2 =  4, x3 = 7)
x
```

```
## x1 x2 x3
##  1  4  7
```

Get the names of a vector:

```
names(x)
```

```
## [1] "x1" "x2" "x3"
```

Named Vectors

The `names()` function can also be used to provide names to a vector:

```r
y <- 1:3
names(y) <- c("y1", "y2", "y3")
```

Subsetting Named Vectors

Vectors can also be subseted by name:

```
y
```

```
## y1 y2 y3
##  1  2  3
```

```
y["y1"]
```

```
## y1
##  1
```

```
y[c("y1", "y3")]
```

```
## y1 y3
##  1  3
```

The paste() and paste0() functions

paste() and paste0() can be useful to generate vector names:

```r
paste("y", 1:length(y), sep = "")
```

```
## [1] "y1" "y2" "y3"
```

```r
paste("name", 1:length(y), sep = "_")
```

```
## [1] "name_1" "name_2" "name_3"
```

```r
paste("year", 1990:1993, sep = "-")
```

```
## [1] "year-1990" "year-1991" "year-1992" "year-1993"
```

```r
paste0("X", 1:5)
```

```
## [1] "X1" "X2" "X3" "X4" "X5"
```