

Introduction to R Programming

Data Frames

Pedro Fonseca

08 May 2020

Before we start

The `ifelse()` function performs vectorized if-then-else statements

```
ifelse(test, x, y)
```

- ▶ `test` can be anything that returns a logical vector or matrix
- ▶ `ifelse()` returns `x` in the positions where the corresponding element of `test` is `TRUE` and `y` in the positions where the corresponding element of `test` is `FALSE`

Before we start

```
vec <- 1:5
```

```
ifelse(vec > 2, 2, 1)
```

```
## [1] 1 1 2 2 2
```

```
ifelse(vec > 2, "large", "small")
```

```
## [1] "small" "small" "large" "large" "large"
```

```
ifelse(vec > 2, vec - 1, vec + 1)
```

```
## [1] 2 3 2 3 4
```

Before we start

You can chain multiple `ifelse()` calls:

```
heights <- c(1.66, 1.88, 1.76, 1.68, 1.7, 1.9)

shirts <- ifelse(heights >= 1.9, "XL",
                 ifelse(heights > 1.8, "L",
                         ifelse(heights > 1.75, "M",
                                "S")
                        )
                 )

shirts

## [1] "S" "L" "M" "S" "S" "XL"
```

Datasets

In a typical dataset:

- ▶ Rows represent observations
- ▶ Columns represent variables

	age	smoker	marital
1	20	TRUE	M
2	38	FALSE	S
3	50	FALSE	W
4	19	TRUE	S
5	15	FALSE	S

Figure 1: A rectangular dataset

What is a data frame?

Data frames are designed specifically for rectangular datasets

- ▶ Why not matrices? We may need different data types for different columns
- ▶ Why not lists? Not very practical nor visually appealing

My first data frame

You can create a data frame by supplying name-vector pairs to `data.frame()`:

```
df <- data.frame(  
  age = c(20, 38, 50, 19),  
  smoker = c(TRUE, FALSE, FALSE, FALSE),  
  marital = factor(c("M", "S", "W", "S"))  
)
```

My first data frame

Let's print df:

```
df
```

```
##    age smoker marital
## 1   20   TRUE        M
## 2   38  FALSE        S
## 3   50  FALSE        W
## 4   19  FALSE        S
```

Notice that data frame rows are numbered by default

Some utilities

```
summary(df)
```

##	age	smoker	marital
##	Min. :19.00	Mode :logical	M:1
##	1st Qu.:19.75	FALSE:3	S:2
##	Median :29.00	TRUE :1	W:1
##	Mean :31.75		
##	3rd Qu.:41.00		
##	Max. :50.00		

Try this:

```
View(df)
```

What is a data frame?

```
str(df)
```

```
## 'data.frame':    4 obs. of  3 variables:  
## $ age      : num  20 38 50 19  
## $ smoker   : logi  TRUE FALSE FALSE FALSE  
## $ marital: Factor w/ 3 levels "M","S","W": 1 2 3 2
```

- ▶ Similar to the structure of a list?
- ▶ Actually, data frames are build on top of lists
- ▶ A data frame is a list of atomic vectors with equal length

What is a data frame?

Data frames have attributes for column names, row.names, and their own class: data.frame:

```
typeof(df)
```

```
## [1] "list"
```

```
class(df)
```

```
## [1] "data.frame"
```

What is a data frame?

Due to their rectangular structure, data frames share the properties of both matrices and lists:

- ▶ A data frame has `rownames()` and `colnames()`. The `names()` of a data frame are the column names
- ▶ A data frame has `nrow()` rows and `ncol()` columns. The `length()` of a data frame is the number of columns

My second data frame

```
name <- c("Anne", "Pete", "Frank", "Julia", "Cath")
age <- c(28, 30, 21, 39, 35)
child <- c(FALSE, TRUE, TRUE, FALSE, TRUE)

df2 <- data.frame(name, age, child)
df2
```

```
##      name age child
## 1  Anne  28 FALSE
## 2  Pete  30  TRUE
## 3 Frank  21  TRUE
## 4 Julia  39 FALSE
## 5  Cath  35  TRUE
```

Unlike `list()`, `data.frame()` keeps the names of the input vectors

Assign names() to a data frame

Data frame columns always have names. If you don't provide them, default names will be assigned:

```
df3 <- data.frame(  
  c(1, 4, 6),  
  c(6, 7, 9)  
)
```

```
df3
```

```
##      c.1..4..6. c.6..7..9.  
## 1             1           6  
## 2             4           7  
## 3             6           9
```

Assign names() to a data frame

```
names(df3) <- c("a-name", "some-other-name")
```

```
df3
```

##	a-name	some-other-name
## 1	1	6
## 2	4	7
## 3	6	9

Row names

`data.frame()` allows you to label row indexes. You can provide a vector of unique values to the `row.names` argument:

```
df4 <- data.frame(  
  age = c(35, 27, 18),  
  hair = c("blond", "brown", "black"),  
  row.names = c("Bob", "Susan", "Sam") # row names!  
)
```

df4

##	age	hair
## Bob	35	blond
## Susan	27	brown
## Sam	18	black

Row names

`rownames()` can be used to get, set or overwrite the row names of an existing data frame:

```
rownames(df4)
```

```
## [1] "Bob"    "Susan" "Sam"
```

```
rownames(df4) <- c("Bob M.", "Susan B.", "Sam L.")
```

```
df4
```

```
##           age  hair
## Bob M.      35 blond
## Susan B.    27 brown
## Sam L.      18 black
```

stringsAsFactors

```
df5 <- data.frame(  
  person = c("John", "Matt", "Mary"),  
  age = c(20, 38, 50),  
  smoker = c(TRUE, FALSE, FALSE),  
  marital = factor(c("M", "S", "S"))  
)  
  
str(df5)
```

```
## 'data.frame':    3 obs. of  4 variables:  
## $ person : Factor w/ 3 levels "John","Mary",...: 1 3 2  
## $ age      : num  20 38 50  
## $ smoker   : logi  TRUE FALSE FALSE  
## $ marital  : Factor w/ 2 levels "M","S": 1 2 2
```

The person variable is a factor! Why?

stringsAsFactors

The `stringsAsFactors` argument of `data.frame` is set to `TRUE` by default.

```
df6 <- data.frame(  
  person = c("John", "Matt", "Mary"),  
  age = c(20, 38, 50),  
  smoker = c(TRUE, FALSE, FALSE),  
  marital = factor(c("M", "S", "S")),  
  stringsAsFactors = FALSE  
)  
str(df6)
```

```
## 'data.frame':    3 obs. of  4 variables:  
## $ person : chr  "John" "Matt" "Mary"  
## $ age : num  20 38 50  
## $ smoker : logi  TRUE FALSE FALSE  
## $ marital: Factor w/ 2 levels "M","S": 1 2 2
```

Testing and coercing

To check if an object is a data frame use `is.data.frame()`:

```
is.data.frame(df6)
```

```
## [1] TRUE
```

You can coerce an object to a data frame with `as.data.frame()`:

```
M <- matrix(1:9, ncol = 3)
as.data.frame(M)
```

```
##   V1 V2 V3
## 1  1  4  7
## 2  2  5  8
## 3  3  6  9
```

Subsetting

Data frames inherit subsetting syntax from both matrices and lists.
You can subset a data frame:

- ▶ Like a 2D structure (behaves like a matrix)
- ▶ Like a 1D structure (behaves like a list)

Subsetting

```
name <- c("Anne", "Pete", "Frank", "Julia", "Cath")
age <- c(28, 30, 21, 39, NA)
child <- c(FALSE, FALSE, TRUE, FALSE, TRUE)

people <- data.frame(name, age, child,
                      stringsAsFactors = FALSE)

people
```

```
##      name age child
## 1  Anne  28 FALSE
## 2  Pete  30 FALSE
## 3 Frank  21  TRUE
## 4 Julia  39 FALSE
## 5  Cath  NA  TRUE
```

Subsetting with matrix syntax

```
people[3, ] # returns a data frame
```

```
##      name age child  
## 3 Frank  21  TRUE
```

```
people[c(1, 3), ] # data frame
```

```
##      name age child  
## 1  Anne  28 FALSE  
## 3 Frank  21  TRUE
```

```
people[3, 2] # vector
```

```
## [1] 21
```

```
people[3, "age"] # vector
```

```
## [1] 21
```

Subsetting with matrix syntax

```
people[, "age"] # vector
```

```
## [1] 28 30 21 39 NA
```

```
people[, 1] # vector
```

```
## [1] "Anne" "Pete" "Frank" "Julia" "Cath"
```

```
people[c(3, 5), c("age", "child")] # data frame
```

```
##   age child
```

```
## 3  21  TRUE
```

```
## 5  NA  TRUE
```


Subsetting with matrix syntax

See the pattern?

- ▶ When you subset a data frame like a 2D object with `[]` you get a vector if only one column is extracted. You get a data frame otherwise.
- ▶ This is a frequent source of bugs when using `[]` in a function, unless you always remember to use `drop = FALSE`.

Subsetting with matrix syntax

```
people[, "age", drop = FALSE] # data frame!
```

```
##    age  
## 1  28  
## 2  30  
## 3  21  
## 4  39  
## 5  NA
```

Setting `drop = FALSE` prevents [from coercing to a lower dimension.

Subsetting with matrix syntax

You can also subset using logical conditions. Show Pete's row:

```
people[people[, "name"] == "Pete", ]
```

```
##   name age child
```

```
## 2 Pete  30 FALSE
```

Subsetting with matrix syntax

Show Julia's age and wheather or not she has children:

```
people[people[, "name"] == "Julia", c("age", "child")]
```

```
##    age child
```

```
## 4   39 FALSE
```

Subsetting with matrix syntax

Excluding Cath, show age and child for people older than 28:

```
people[people[, "name"] != "Cath" &  
        people[, "age"] > 28,  
        c("name", "child")]
```

```
##      name child  
## 2    Pete FALSE  
## 4   Julia FALSE
```

Subsetting with matrix syntax

Filter out observations with missing age:

```
people[!is.na(people[, "age"]), ]
```

```
##      name age child  
## 1  Anne  28 FALSE  
## 2  Pete  30 FALSE  
## 3 Frank  21  TRUE  
## 4 Julia  39 FALSE
```

Subsetting with matrix syntax

Who's age is missing?

```
people[is.na(people[, "age"]), "name"]
```

```
## [1] "Cath"
```

Subsetting with matrix syntax

Show name and child for the people who are older than 21:

```
people[people[, "age"] > 21, c("name", "child")]
```

```
##      name child
## 1    Anne FALSE
## 2    Pete FALSE
## 4   Julia FALSE
## NA  <NA>    NA
```

Why do we have a row of NAs?

Subsetting with matrix syntax

```
table(is.na(people$age))
```

```
##  
## FALSE  TRUE  
##      4      1
```

Solution:

```
people[!is.na(people[, "age"]) & people[, "age"] > 21,  
       c("name", "child")]
```

```
##      name child  
## 1  Anne FALSE  
## 2  Pete FALSE  
## 4  Julia FALSE
```

Subsetting with matrix syntax

Another solution:

```
people_with_age <- people[!is.na(people[, "age"]), ]  
  
people_with_age[people_with_age[, "age"] > 21,  
  c("name", "child")]
```

```
##      name child  
## 1  Anne FALSE  
## 2   Pete FALSE  
## 4  Julia FALSE
```

Subsetting with matrix syntax

If available, row names can be used for subsetting:

```
df4
```

```
##           age  hair  
## Bob M.     35 blond  
## Susan B.   27 brown  
## Sam L.     18 black
```

```
df4["Bob M.", ]
```

```
##           age  hair  
## Bob M.     35 blond
```

Subsetting with matrix syntax

```
df4
```

```
##           age  hair  
## Bob M.     35 blond  
## Susan B.   27 brown  
## Sam L.     18 black
```

```
df4[c("Bob M.", "Sam L."), "age"]
```

```
## [1] 35 18
```

Subsetting with list syntax

```
people
```

```
##      name age child
## 1  Anne  28 FALSE
## 2  Pete  30 FALSE
## 3 Frank  21  TRUE
## 4 Julia  39 FALSE
## 5  Cath  NA  TRUE
```

Subsetting with list syntax

```
people[2] # data frame
```

```
##    age  
## 1  28  
## 2  30  
## 3  21  
## 4  39  
## 5  NA
```

```
people["age"] # data frame
```

```
##    age  
## 1  28  
## 2  30  
## 3  21  
## 4  39  
## 5  NA
```

Subsetting with list syntax

```
people[c(1, 3)] # data frame
```

```
##      name child  
## 1  Anne FALSE  
## 2   Pete FALSE  
## 3  Frank  TRUE  
## 4  Julia FALSE  
## 5   Cath  TRUE
```

```
people[c("name", "child")] # data frame
```

```
##      name child  
## 1  Anne FALSE  
## 2   Pete FALSE  
## 3  Frank  TRUE  
## 4  Julia FALSE  
## 5   Cath  TRUE
```

Subsetting with list syntax

```
people["age"] # data frame
```

```
##    age  
## 1   28  
## 2   30  
## 3   21  
## 4   39  
## 5  NA
```

```
people[["age"]] # vector
```

```
## [1] 28 30 21 39 NA
```

```
people$age # vector
```

```
## [1] 28 30 21 39 NA
```


Subsetting with list syntax

```
people[["age"]][1]
```

```
## [1] 28
```

```
people[["age"]][c(1, 4)]
```

```
## [1] 28 39
```

```
people$age[1]
```

```
## [1] 28
```

```
people$age[c(1, 4)]
```

```
## [1] 28 39
```

Subsetting with list syntax

```
people[[2]] # same as people[["age"]]
```

```
## [1] 28 30 21 39 NA
```

```
people[[2]][1]
```

```
## [1] 28
```

```
people[[2]][c(1, 4)]
```

```
## [1] 28 39
```

Subsetting with list syntax

Since data frames are lists of vectors, when you subset a data frame like a 1D object:

- ▶ with `[` you get a data frame.
- ▶ with `[[` or `$` you get a vector.

Partial matching with \$

When you attempt to extract a column from `df` with `df$name` and there is no column `name`, a data frame will instead select any variable that starts with `name`. If no variable starts with `name` (or if more than one do), `df$name` will return `NULL`.

```
names(df)
```

```
## [1] "age"      "smoker"   "marital"
```

```
df$a
```

```
## [1] 20 38 50 19
```

By default, `[[` does not do partial matching. Why? Run `?"["` and check what is the default value of the `exact` argument.

`df$name` is equivalent to `df[["name", exact = FALSE]]`

Adding new columns

So far our data frame `people` has 5 rows and 3 columns.

```
str(people)
```

```
## 'data.frame':    5 obs. of  3 variables:
## $ name : chr  "Anne" "Pete" "Frank" "Julia" ...
## $ age  : num  28 30 21 39 NA
## $ child: logi  FALSE FALSE TRUE FALSE TRUE
```

Adding columns with [

Let's add a column with people's height:

```
heights <- c(175, 170, 166, 182, 172)
```

```
people["height"] <- heights
```

```
str(people)
```

```
## 'data.frame':    5 obs. of  4 variables:
```

```
## $ name   : chr  "Anne" "Pete" "Frank" "Julia" ...
```

```
## $ age    : num  28 30 21 39 NA
```

```
## $ child  : logi  FALSE FALSE TRUE FALSE TRUE
```

```
## $ height: num  175 170 166 182 172
```

Adding columns with [

We can use the position of the new column instead of the name:

```
has_dog <- factor(c("yes", "no", "yes", "no", "no"))  
people[5] <- has_dog
```

```
str(people)
```

```
## 'data.frame':    5 obs. of  5 variables:  
## $ name   : chr  "Anne" "Pete" "Frank" "Julia" ...  
## $ age    : num  28 30 21 39 NA  
## $ child  : logi  FALSE FALSE TRUE FALSE TRUE  
## $ height : num  175 170 166 182 172  
## $ V5     : Factor w/ 2 levels "no","yes": 2 1 2 1 1
```

Adding columns with [

What is the consequence of using the position of the new column instead of the name?

► Hint: look at the names of the new columns.

Lets rename V5:

```
names(people)[5] <- "dog"
```

```
people
```

##		name	age	child	height	dog
## 1	Anne	28	FALSE	175	yes	
## 2	Pete	30	FALSE	170	no	
## 3	Frank	21	TRUE	166	yes	
## 4	Julia	39	FALSE	182	no	
## 5	Cath	NA	TRUE	172	no	

Adding columns with \$

Let's add a column with people's weight using \$:

```
weight <- c(86, 63, 68, 55, 56)
people$weight <- weight

str(people)
```

```
## 'data.frame':    5 obs. of  6 variables:
##  $ name   : chr   "Anne" "Pete" "Frank" "Julia" ...
##  $ age    : num   28 30 21 39 NA
##  $ child  : logi  FALSE FALSE TRUE FALSE TRUE
##  $ height : num   175 170 166 182 172
##  $ dog    : Factor w/ 2 levels "no","yes": 2 1 2 1 1
##  $ weight : num   86 63 68 55 56
```

With \$ we can only add new columns by name.

Adding columns with cbind()

Now let's add a column with the people's income:

```
income <- c(1200, 750, 1660, 1280, 890)
people <- cbind(people, income)
```

```
str(people)
```

```
## 'data.frame':    5 obs. of  7 variables:
## $ name   : chr   "Anne" "Pete" "Frank" "Julia" ...
## $ age    : num   28 30 21 39 NA
## $ child  : logi   FALSE FALSE TRUE FALSE TRUE
## $ height: num   175 170 166 182 172
## $ dog     : Factor w/ 2 levels "no","yes": 2 1 2 1 1
## $ weight: num    86 63 68 55 56
## $ income: num   1200 750 1660 1280 890
```

Dropping columns with list syntax

Let's drop the dog column. By name:

```
people$dog <- NULL  
# Or: people["dog"] <- NULL
```

```
people
```

	##	name	age	child	height	weight	income
##	1	Anne	28	FALSE	175	86	1200
##	2	Pete	30	FALSE	170	63	750
##	3	Frank	21	TRUE	166	68	1660
##	4	Julia	39	FALSE	182	55	1280
##	5	Cath	NA	TRUE	172	56	890

Dropping columns with list syntax

Now let's drop income. By position:

```
people[6] <- NULL  
# or: people <- people[-6]
```

```
people
```

	name	age	child	height	weight
## 1	Anne	28	FALSE	175	86
## 2	Pete	30	FALSE	170	63
## 3	Frank	21	TRUE	166	68
## 4	Julia	39	FALSE	182	55
## 5	Cath	NA	TRUE	172	56

Reordering columns

To reorder columns, just select them in the desired order. For example, let's swap the height and weight columns:

```
people <- people[, c("name", "age", "child", "weight", "height")]
```

```
people
```

	name	age	child	weight	height
## 1	Anne	28	FALSE	86	175
## 2	Pete	30	FALSE	63	170
## 3	Frank	21	TRUE	68	166
## 4	Julia	39	FALSE	55	182
## 5	Cath	NA	TRUE	56	172

Modifying columns

We can overwrite existing columns. Lets convert height from centimeters to meters:

```
people$height <- people$height/100
```

```
people
```

	##	name	age	child	weight	height
##	1	Anne	28	FALSE	86	1.75
##	2	Pete	30	FALSE	63	1.70
##	3	Frank	21	TRUE	68	1.66
##	4	Julia	39	FALSE	55	1.82
##	5	Cath	NA	TRUE	56	1.72

Modifying columns

In the previous example, the following commands would also have worked:

```
people["height"] <- people["height"]/100  
people[5] <- people[5]/100
```

Calculations with existing columns

We can create new columns using the values of existing columns. Let's use `weight` and `height` to calculate the body mass index:

```
people$bmi <- round(  
  people$weight/((people$height)^2)  
)  
people
```

	##	name	age	child	weight	height	bmi
##	1	Anne	28	FALSE	86	1.75	28
##	2	Pete	30	FALSE	63	1.70	22
##	3	Frank	21	TRUE	68	1.66	25
##	4	Julia	39	FALSE	55	1.82	17
##	5	Cath	NA	TRUE	56	1.72	19

Calculations with existing columns

Now let's use the body mass index to calculate the weight status:

```
people$weight_status <- ifelse(  
  people$bmi < 18.5, "underweight",  
    ifelse(people$bmi < 24.9, "normal",  
      ifelse(people$bmi < 29.9, "overweight",  
        "obese")  
    )  
  )  
)
```

Calculations with existing columns

```
people
```

##	name	age	child	weight	height	bmi	weight_status
## 1	Anne	28	FALSE	86	1.75	28	overweight
## 2	Pete	30	FALSE	63	1.70	22	normal
## 3	Frank	21	TRUE	68	1.66	25	overweight
## 4	Julia	39	FALSE	55	1.82	17	underweight
## 5	Cath	NA	TRUE	56	1.72	19	normal

Dropping columns using matrix syntax

Let's drop the bmi and weight_status columns By name:

```
people[, c("bmi", "weight_status")] <- NULL
```

Now let's drop weight:

```
people[, 4] <- NULL  
# or people <- people[, -4]
```

Add a row

```
tom <- data.frame("Tom", 37, FALSE, 1.83)
```

```
rbind(people, tom)
```

```
## Error in match.names(clabs, names(xi)) : names do  
## not match previous names
```

Add a row

The default names do not match the names of the people data frame:

```
names(tom)
```

```
## [1] "X.Tom." "X37"    "FALSE." "X1.83"
```

```
names(people)
```

```
## [1] "name"    "age"     "child"   "height"
```

Add a row

```
tom <- data.frame(name = "Tom", age = 37,  
                  child = FALSE, height = 1.83)  
rbind(people, tom)
```

```
##   name age child height  
## 1  Anne  28 FALSE   1.75  
## 2  Pete  30 FALSE   1.70  
## 3 Frank  21  TRUE   1.66  
## 4 Julia  39 FALSE   1.82  
## 5  Cath  NA  TRUE   1.72  
## 6   Tom  37 FALSE   1.83
```

Add a row

You can also add rows using vectors, as long as they have appropriate names and the correct length:

```
new_obs <- t(c(name = "Peter", age = 30,  
               child = FALSE, height = 1.86))
```

```
rbind(people, new_obs)
```

```
##   name  age child height  
## 1 Anne   28 FALSE   1.75  
## 2 Pete   30 FALSE    1.7  
## 3 Frank  21  TRUE   1.66  
## 4 Julia  39 FALSE   1.82  
## 5 Cath <NA>  TRUE   1.72  
## 6 Peter  30 FALSE   1.86
```

Combining matrix and list syntax

Show the name and age of the people without children:

```
people[people$child == FALSE, c("name", "age")]
```

```
##      name age
## 1  Anne  28
## 2  Pete  30
## 4 Julia  39
```


Combining matrix and list syntax

Show the name and age of the people who are taller than 1.70:

```
people[people$height > 1.70, c("name", "age")]
```

```
##      name age
## 1  Anne  28
## 4 Julia  39
## 5  Cath  NA
```

Combining matrix and list syntax

Show the name and age of the people without children who are taller than 1.70:

```
people[people$child == FALSE & people$height > 1.70,  
       c("name", "age")]
```

```
##      name age  
## 1  Anne  28  
## 4 Julia  39
```

Combining matrix and list syntax

Show the name, age and height of the people with children who are taller than 1.70:

```
people[people$child == TRUE & people$height > 1.70,  
       c("name", "age", "height")]
```

```
##   name age height
```

```
## 5 Cath  NA   1.72
```

Sorting

```
people$age
```

```
## [1] 28 30 21 39 NA
```

```
sort(people$age)
```

```
## [1] 21 28 30 39
```

```
ranks <- order(people$age)  
ranks
```

```
## [1] 3 1 2 4 5
```

```
people$age[ranks]
```

```
## [1] 21 28 30 39 NA
```

Sorting

Sort the rows by ascending age:

```
people[ranks, ]
```

```
##      name age child height
## 3 Frank  21  TRUE   1.66
## 1  Anne  28 FALSE   1.75
## 2  Pete  30 FALSE   1.70
## 4 Julia  39 FALSE   1.82
## 5  Cath  NA  TRUE   1.72
```

Sorting

Sort the rows by descending age:

```
people[order(-people$age), ]
```

```
##      name age child height
## 4 Julia  39 FALSE   1.82
## 2 Pete   30 FALSE   1.70
## 1 Anne   28 FALSE   1.75
## 3 Frank  21  TRUE   1.66
## 5 Cath   NA  TRUE   1.72
```

Sorting

Sort the rows by descending height:

```
people[order(people$height, decreasing = TRUE), ]
```

```
##      name age child height
## 4 Julia  39 FALSE   1.82
## 1 Anne  28 FALSE   1.75
## 5 Cath  NA  TRUE   1.72
## 2 Pete  30 FALSE   1.70
## 3 Frank 21  TRUE   1.66
```

Sorting

Sort the rows by name:

```
people[order(people$name), ]
```

```
##      name age child height
## 1  Anne  28 FALSE   1.75
## 5  Cath  NA  TRUE   1.72
## 3 Frank  21  TRUE   1.66
## 4 Julia  39 FALSE   1.82
## 2  Pete  30 FALSE   1.70
```