

Introduction to R Programming

The Apply Family

Pedro Fonseca

30 April 2022

Built in datasets

R provides many built in datasets. For a complete list see:

```
library(help = "datasets")
```

In this lecture we will use the following datasets:

- ▶ USArrests
- ▶ airquality

For information about a specific dataset see, for example:

```
?USArrests
```

The USArrests dataset

This data set contains:

- ▶ Arrests per 100.000 residents for assault, murder and rape in each of the 50 USA states in 1973.
- ▶ There is also a column with the percentage of population living in urban areas.

The USArrests dataset

```
head(USArrests)
```

##		Murder	Assault	UrbanPop	Rape
##	Alabama	13.2	236	58	21.2
##	Alaska	10.0	263	48	44.5
##	Arizona	8.1	294	80	31.0
##	Arkansas	8.8	190	50	19.5
##	California	9.0	276	91	40.6
##	Colorado	7.9	204	78	38.7

The USArrests dataset

Let's do some calculations with the arrests of four states:

```
USArrests_short <- USArrests[1:4, -3]
```

```
USArrests_short
```

##		Murder	Assault	Rape
##	Alabama	13.2	236	21.2
##	Alaska	10.0	263	44.5
##	Arizona	8.1	294	31.0
##	Arkansas	8.8	190	19.5

Row and column sums

```
colSums(USArrests_short)
```

```
## Murder Assault Rape  
## 40.1 983.0 116.2
```

```
rowSums(USArrests_short)
```

```
## Alabama Alaska Arizona Arkansas  
## 270.4 317.5 333.1 218.3
```

Row and column means

```
colMeans(USArrests_short)
```

```
## Murder Assault Rape  
## 10.025 245.750 29.050
```

```
rowMeans(USArrests_short)
```

```
## Alabama Alaska Arizona Arkansas  
## 90.13333 105.83333 111.03333 72.76667
```

The `apply()` function

To collapse data frames across rows or columns using functions other than the sum and the mean we can use `apply()`:

```
apply(X, MARGIN, FUN, ...)
```

- ▶ `X` is a data frame
- ▶ `MARGIN = 1` for rows, `MARGIN = 2` for columns
- ▶ `FUN` is a function
- ▶ `...` are optional arguments to pass to `FUN`

The apply() function

Apply functions over the columns of USArrests_short:

```
apply(USArrests_short, 2, mean)
```

```
## Murder Assault Rape  
## 10.025 245.750 29.050
```

```
apply(USArrests_short, 2, median)
```

```
## Murder Assault Rape  
## 9.4 249.5 26.1
```

```
apply(USArrests_short, 2, sd)
```

```
## Murder Assault Rape  
## 2.257395 44.078528 11.479402
```

The apply() function

Apply functions the rows of USArrests_short:

```
apply(USArrests_short, 1, max)
```

```
## Alabama Alaska Arizona Arkansas  
##      236      263      294      190
```

```
apply(USArrests_short, 1, min)
```

```
## Alabama Alaska Arizona Arkansas  
##      13.2      10.0       8.1       8.8
```

```
apply(USArrests_short, 1, var)
```

```
## Alabama Alaska Arizona Arkansas  
## 15973.81 18823.58 25238.70 10336.36
```

apply() vs for loop

Loop over the columns of USArrests_short:

```
res <- vector()

for(i in 1:ncol(USArrests_short)){

  res[i] <- mean(USArrests_short[[i]], na.rm = TRUE)
  names(res)[i] <- names(USArrests_short)[i]
}
```

res

```
## Murder Assault Rape
## 10.025 245.750 29.050
```

apply() vs for loop

Loop over the rows of USArrests_short:

```
res <- vector()

for(j in 1:nrow(USArrests_short)){

  res[j] <- max(USArrests_short[j, ], na.rm = TRUE)
  names(res)[j] <- rownames(USArrests_short)[j]
}
```

res

```
## Alabama   Alaska   Arizona Arkansas
##      236      263      294      190
```

apply() with ...

Now let's see an example that requires using dot-dot-dot (...). Suppose you're working with the `airquality` the data frame:

```
head(airquality)
```

##	Ozone	Solar.R	Wind	Temp	Month	Day
## 1	41	190	7.4	67	5	1
## 2	36	118	8.0	72	5	2
## 3	12	149	12.6	74	5	3
## 4	18	313	11.5	62	5	4
## 5	NA	NA	14.3	56	5	5
## 6	28	NA	14.9	66	5	6

`apply()` with ...

Try to use `apply()` to compute the means of the first four columns:

```
apply(airquality[, 1:4], 2, mean)
```

```
##      Ozone      Solar.R      Wind      Temp  
##         NA         NA  9.957516 77.882353
```

We get NA for the first two columns. Why?

apply() with ...

Problem: Some columns have NAs.

```
sum(is.na(airquality$Ozone))
```

```
## [1] 37
```

```
sum(is.na(airquality$Solar.R))
```

```
## [1] 7
```

`apply()` with ...

Solution: use ... to pass the `na.rm` argument to `mean()`:

```
apply(airquality[, 1:4], 2, mean, na.rm = TRUE)
```

##	Ozone	Solar.R	Wind	Temp
##	42.129310	185.931507	9.957516	77.882353

apply() user-defined functions

Checking for NAs one column at a time is not be feasible when we have many columns. How can we check how many NAs there there in each column?

```
count_nas <- function(x){sum(is.na(x))}
```

```
apply(airquality, 2, count_nas)
```

##	Ozone	Solar.R	Wind	Temp	Month	Day
##	37	7	0	0	0	0

The `lapply()` function

Syntax: `lapply(X, FUN, ...)`

`lapply()` is similar to `apply()` but:

- ▶ `X` is a vector (atomic or list)
- ▶ There is no `MARGINS` argument
- ▶ Always returns a list

`lapply()` returns a list of the same length as `X`, each element of which is the result of applying `FUN` to the corresponding element of `X`.

The lapply() function

```
lapply(c(1:3), log, base = 10)
```

```
## [[1]]
```

```
## [1] 0
```

```
##
```

```
## [[2]]
```

```
## [1] 0.30103
```

```
##
```

```
## [[3]]
```

```
## [1] 0.4771213
```

The lapply() function

```
A <- matrix(1:10, ncol = 5)
B <- matrix(c(1, 5, 7, -1), ncol = 4)
C <- matrix(letters[1:4], ncol = 2)

my_list <- list(A, B, C)
```

The lapply() function

```
my_list
```

```
## [[1]]  
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]    1    3    5    7    9  
## [2,]    2    4    6    8   10  
##  
## [[2]]  
##      [,1] [,2] [,3] [,4]  
## [1,]    1    5    7   -1  
##  
## [[3]]  
##      [,1] [,2]  
## [1,] "a"  "c"  
## [2,] "b"  "d"
```

The `lapply()` function

Every element of `my_list`, except the last, contain a numerical matrix. Sum the elements of each of those matrices:

```
lapply(my_list[-3], sum)
```

```
## [[1]]
```

```
## [1] 55
```

```
##
```

```
## [[2]]
```

```
## [1] 12
```

The `lapply()` function

Extract the element in position (1, 2) from each matrix:

```
lapply(my_list, "[", 1, 2)
```

```
## [[1]]
```

```
## [1] 3
```

```
##
```

```
## [[2]]
```

```
## [1] 5
```

```
##
```

```
## [[3]]
```

```
## [1] "c"
```

The `lapply()` function

Extract the first row from each matrix:

```
lapply(my_list,"[, 1 , )
```

```
## [[1]]
```

```
## [1] 1 3 5 7 9
```

```
##
```

```
## [[2]]
```

```
## [1] 1 5 7 -1
```

```
##
```

```
## [[3]]
```

```
## [1] "a" "c"
```


The `lapply()` function

Extract the 2nd column from each matrix:

```
lapply(my_list,"[,", 2)
```

```
## [[1]]
```

```
## [1] 3 4
```

```
##
```

```
## [[2]]
```

```
## [1] 5
```

```
##
```

```
## [[3]]
```

```
## [1] "c" "d"
```

lapply() vs for loop

Extract the 2nd column from each matrix:

```
res <- vector(mode = "list")

for(i in seq_along(my_list)){
  res[[i]] <- my_list[[i]][, 2]
}

res
```

```
## [[1]]
## [1] 3 4
##
## [[2]]
## [1] 5
##
## [[3]]
## [1] "c" "d"
```

The `sapply()` function

The `sapply()` function:

- ▶ Works like `lapply()`, but simplifies the output to the most elementary data structure that is possible.
- ▶ Returns vectors or matrices.

The sapply() function

```
sapply(my_list[-3], sum)
```

```
## [1] 55 12
```

```
sapply(my_list, "[", 1, 2)
```

```
## [1] "3" "5" "c"
```

The sapply() function

```
sapply(airquality, function(x) sum(is.na(x)))
```

##	Ozone	Solar.R	Wind	Temp	Month	Day
##	37	7	0	0	0	0

The `apply()` function

```
set.seed(123)

our_list <- list(
  w = 1:6,
  x = sample(1:5, 4, replace = TRUE),
  y = matrix(sample(1:100, 9), nrow = 3),
  z = sample(1:10, 3, replace = TRUE)
)

str(our_list)

## List of 4
## $ w: int [1:6] 1 2 3 4 5 6
## $ x: int [1:4] 3 3 2 2
## $ y: int [1:3, 1:3] 43 14 25 90 91 69 96 57 92
## $ z: int [1:3] 9 9 3
```

The `sapply()` function

```
sapply(our_list, max)
```

```
##  w  x  y  z  
##  6  3 96  9
```

```
sapply(our_list, min)
```

```
##  w  x  y  z  
##  1  2 14  3
```

```
sapply(our_list, class)
```

```
## $w  
## [1] "integer"  
##  
## $x  
## [1] "integer"  
##  
## $y  
## [1] "matrix" "array"
```

The `sapply()` function

How many numbers are there inside each element of `our_list`?

```
sapply(our_list, length)
```

```
## w x y z
```

```
## 6 4 9 3
```


The `sapply()` function

```
set.seed(123)
```

```
our_list_2 <- list(  
  w = 1:6,  
  x = sample(1:5, 4, replace = TRUE),  
  y = airquality  
)
```

```
sapply(our_list_2, class)
```

```
##           w           x           y  
## "integer" "integer" "data.frame"
```

```
dim(airquality)
```

```
## [1] 153   6
```

How many numbers are there inside each element of `our_list_2`?

The `sapply()` function

The length of a data frame is the number of columns, and hence `sapply(our_list_2, length)` won't do the trick.

```
our_fun <- function(x){  
  if(class(x) == "data.frame"){  
    nrow(x) * ncol(x)  
  }else{  
    length(x)  
  }  
}
```

```
sapply(our_list_2, our_fun)
```

```
##      w      x      y  
##      6      4 918
```

apply() vs for loop

The same but with a for loop:

```
res <- vector()

for(i in seq_along(our_list_2)){

  res[i] <- our_fun(our_list_2[[i]])
  names(res)[i] <- names(our_list_2)[i]
}
```

```
res
```

```
##      w      x      y
##      6      4 918
```

The `mapply()` function

`mapply()` is a generalization of `sapply()`. It applies a multivariate function over multiple vectors of arguments.

The `mapply()` function

Suppose we want 3 samples of different sizes from a $\text{Normal}(0, 1)$ distribution:

```
set.seed(123)
```

```
rnorm(n = 1)
```

```
## [1] -0.5604756
```

```
rnorm(n = 2)
```

```
## [1] -0.2301775  1.5587083
```

```
rnorm(n = 3)
```

```
## [1] 0.07050839 0.12928774 1.71506499
```

The mapply() function

The same result can be obtained more compactly with mapply():

```
set.seed(123)
```

```
sample_size <- 1:3
```

```
mapply(FUN = rnorm, n = sample_size)
```

```
## [[1]]
```

```
## [1] -0.5604756
```

```
##
```

```
## [[2]]
```

```
## [1] -0.2301775  1.5587083
```

```
##
```

```
## [[3]]
```

```
## [1] 0.07050839 0.12928774 1.71506499
```

The `mapply()` function

Since we only iterated over one vector, we could have used `sapply()`:

```
set.seed(123)

sample_size <- 1:3
sapply(FUN = rnorm, X = sample_size)
```

```
## [[1]]
## [1] -0.5604756
##
## [[2]]
## [1] -0.2301775  1.5587083
##
## [[3]]
## [1] 0.07050839 0.12928774 1.71506499
```

The `mapply()` function

But what if we want to sample from normal distributions with different means, while still having samples of different sizes?

```
set.seed(123)
```

```
rnorm(n = 1, mean = 5)
```

```
## [1] 4.439524
```

```
rnorm(n = 2, mean = 10)
```

```
## [1] 9.769823 11.558708
```

```
rnorm(n = 3, mean = -3)
```

```
## [1] -2.929492 -2.870712 -1.284935
```


The mapply() function

In this case we need to iterate over two vectors, one for sample sizes and one for means:

```
set.seed(123)

sample_size <- 1:3
mu <- c(5, 10, -3)
mapply(rnorm, n = sample_size, mean = mu)

## [[1]]
## [1] 4.439524
##
## [[2]]
## [1] 9.769823 11.558708
##
## [[3]]
## [1] -2.929492 -2.870712 -1.284935
```

The mapply() function

Now suppose we also want each sample to have a different standard deviation:

```
set.seed(123)
```

```
rnorm(n = 1, mean = 5, sd = 1)
```

```
## [1] 4.439524
```

```
rnorm(n = 2, mean = 10, sd = 3)
```

```
## [1] 9.309468 14.676125
```

```
rnorm(n = 3, mean = -3, sd = 5)
```

```
## [1] -2.647458 -2.353561 5.575325
```

The mapply() function

```
set.seed(123)

sample_size <- 1:3
mu <- c(5, 10, -3)
sigma <- c(1, 3, 5)

mapply(rnorm, mean = mu, sd = sigma, n = sample_size)

## [[1]]
## [1] 4.439524
##
## [[2]]
## [1] 9.309468 14.676125
##
## [[3]]
## [1] -2.647458 -2.353561 5.575325
```

The mapply() function

Now suppose we wanted our results with two decimal places only:

```
set.seed(123)
```

```
results <- mapply(rnorm, mean = mu, sd = sigma,  
                  n = sample_size)
```

```
sapply(results, FUN = round, 2)
```

```
## [[1]]
```

```
## [1] 4.44
```

```
##
```

```
## [[2]]
```

```
## [1] 9.31 14.68
```

```
##
```

```
## [[3]]
```

```
## [1] -2.65 -2.35 5.58
```

mapply() vs for loop

```
set.seed(123)

sample_size <- 1:3
mu <- c(5, 10, -3)
sigma <- c(1, 3, 5)

res <- vector(mode = "list")

for (i in 1:3) {
  res[[i]] <- round(rnorm(mean = mu[i],
                          sd = sigma[i],
                          n = sample_size[i]),
                    2)
}
```

mapply() vs for loop

```
res
```

```
## [[1]]
```

```
## [1] 4.44
```

```
##
```

```
## [[2]]
```

```
## [1] 9.31 14.68
```

```
##
```

```
## [[3]]
```

```
## [1] -2.65 -2.35 5.58
```