

Introduction to R Programming

Data Visualitazion with ggplot2

Pedro Fonseca

19 Abril 2020

ggplot2:

Build a data MASTERPIECE



Introduction

- ▶ In lecture you will learn how to visualise data with the ggplot2 package.
- ▶ ggplot2 is one of the most elegant and versatile systems for making graphs.
- ▶ ggplot2 implements the grammar of graphics (Wickham 2010), a coherent system for describing and building graphs.

Preliminars

```
library(ggplot2)
```

- ▶ To access the datasets, help pages, and functions that we will use, load the ggplot2 package.

The mpg dataset

The mpg dataset contains information about 38 models of cars. Among the variables in mpg are:

- ▶ displ: engine size, in litres
- ▶ cyl: number of cylinders
- ▶ cty: city miles per gallon
- ▶ hwy: fuel efficiency on the highway, in miles per gallon (mpg).
- ▶ class: type of car

```
View(mpg)
```

For additional information see `?mpg`.

My first ggplot()

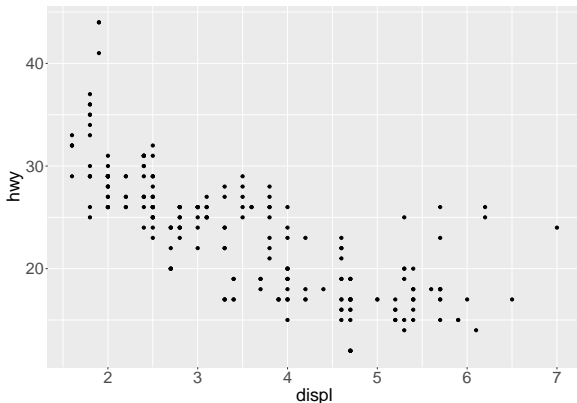
Let's build a graph to answer the following questions:

- ▶ Do cars with big engines use more fuel than cars with small engines?
- ▶ What does the relationship between engine size and fuel efficiency look like? Is it positive? Negative? Linear? Nonlinear?

We start by plotting engine size (`displ`) versus fuel efficiency (`hwy`).

My first ggplot()

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))
```



My first ggplot()

- ▶ The plot shows a negative relationship between engine size and fuel efficiency.
- ▶ In other words: on average, cars with big engines use more fuel.

A simple ggplot() template

```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(mapping = aes(<MAPPINGS>))
```

- ▶ The first argument of ggplot() is a dataset
- ▶ You complete your graph by adding layers to ggplot()
- ▶ geom_point() adds a layer of points, creating scatterplot.
- ▶ ggplot2 has many geom functions that each add a different type of layer to a plot.

A simple `ggplot()` template

- ▶ Each `geom` function takes a `mapping` argument.
- ▶ The `mapping` argument is always paired with `aes()`.
- ▶ The `x` and `y` arguments of `aes()` specify which variables to map to the `x` and `y` axes.

Geometrical objects

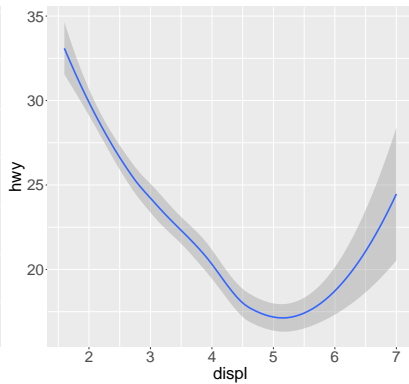
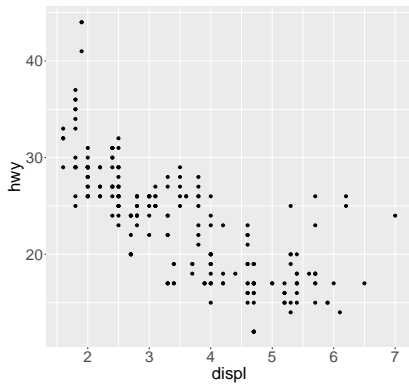
- ▶ A **geom** is the geometrical object that a plot uses to represent data (e.g. points, bars, lines. . .).
- ▶ People often describe plots by the type of **geom** that it uses:
 - ▶ Bar charts use **bar** geoms
 - ▶ Line charts use **line** geoms
 - ▶ Boxplots use **boxplot** geom
 - ▶ . . .
- ▶ Scatterplots break the trend: they use the **point** geom.
- ▶ In ggplot2, you add geoms to a plot with **geom functions**.
- ▶ Different geom functions add different geoms to the plot.

Geometrical objects

Compare the next two plots. How are they similar?

```
left <- ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))  
  
right <- ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))
```

Geometrical objects



Geometrical objects

Both plots:

- ▶ use the same data
- ▶ have the same `x` variable
- ▶ have the same `y` variable

But they are not identical:

- ▶ Each plot uses a different **geom**.
- ▶ The plot on the left uses the **point** geom.
- ▶ The plot on the right uses the **smooth** geom, a smooth line fitted to the data.

The ggplot2 cheatsheet

- ▶ ggplot2 provides over 40 geoms
- ▶ The best way to get a comprehensive overview is the [ggplot2 cheatsheet](#).

Aesthetics

- ▶ **Aesthetics** are visual properties of the **geoms**.
- ▶ **Aesthetics** include things like the size, shape, and color of the points in a scatterplot.
- ▶ `aes()` builds **aesthetic mappings** that define how variables in the dataset are mapped to aesthetics of the geoms.
- ▶ To map an aesthetic to a variable, associate the name of the aesthetic to the name of the variable inside `aes()`.
- ▶ Every geom function needs an aesthetic mapping, but not every aesthetic works with every geom, for example:
 - ▶ You can set the shape of a point, but not of a line.
 - ▶ You can set the linetype of a line, but not of a point.

My first ggplot() revisited

One group of points (highlighted in red) seems to fall outside the linear trend. How can we explain these cars?

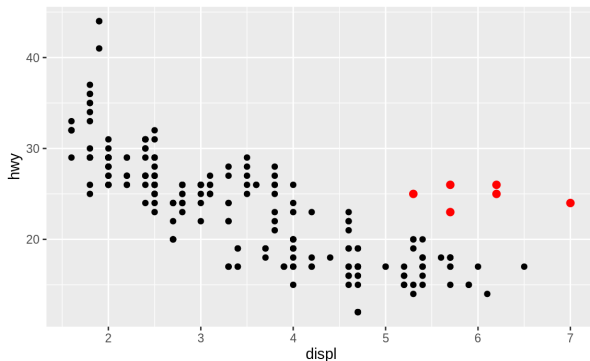


Figure 1: Some cars have a higher mileage than we might expect

My first ggplot() revisited

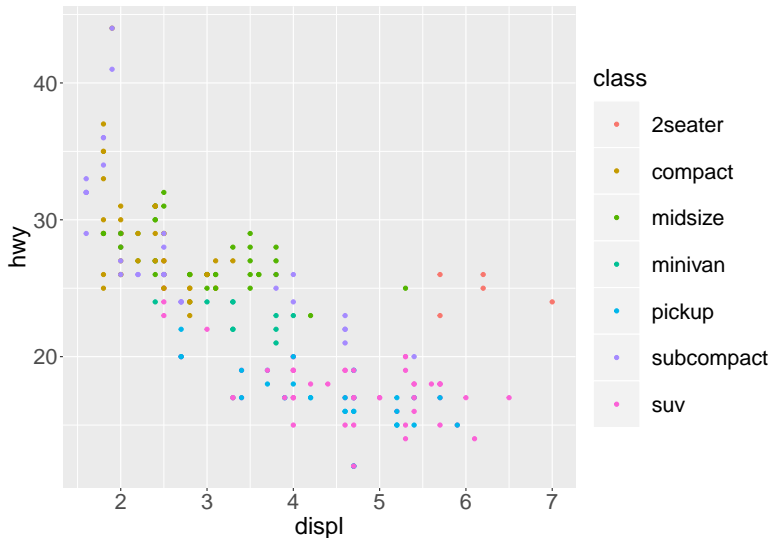
- ▶ The highlighted cars may have some common characteristic with respect to some other variable (e.g. they might all be hybrids).
- ▶ Lets try `class`, a variable that classifies cars into groups such as compact, midsize, and SUV.
- ▶ We can add a third variable to a two dimensional scatterplot by mapping it to an aesthetic.

The color aesthetic

For example, we can map the color of the points to the class variable, so that the graph reveals the class of each car:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy, color = class))
```

The color aesthetic



The color aesthetic

Now we know that all but one of the highlighted cars are two-seater cars! Let's find their model and manufacturer:

```
subset(mpg, class == "2seater",  
       select = c("manufacturer", "model", "year"))
```

```
## # A tibble: 5 x 3  
##   manufacturer model      year  
##   <chr>         <chr>    <int>  
## 1 chevrolet    corvette  1999  
## 2 chevrolet    corvette  1999  
## 3 chevrolet    corvette  2008  
## 4 chevrolet    corvette  2008  
## 5 chevrolet    corvette  2008
```

The color aesthetic

- ▶ These cars are, in fact, sports cars!
- ▶ Sports cars have large engines like SUVs and pickup trucks, but small bodies like midsize and compact cars, which improves their gas mileage.




Figure 2: 2008 Chevrolet Corvette

The size aesthetic

Now let's map class to the size aesthetic:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy, size = class))
```

The size aesthetic



The size aesthetic

- ▶ Each level of the class variable is assigned to a different size.
- ▶ This plot, however, came with a warning message:

```
#> Using size for a discrete variable is not advised.
```

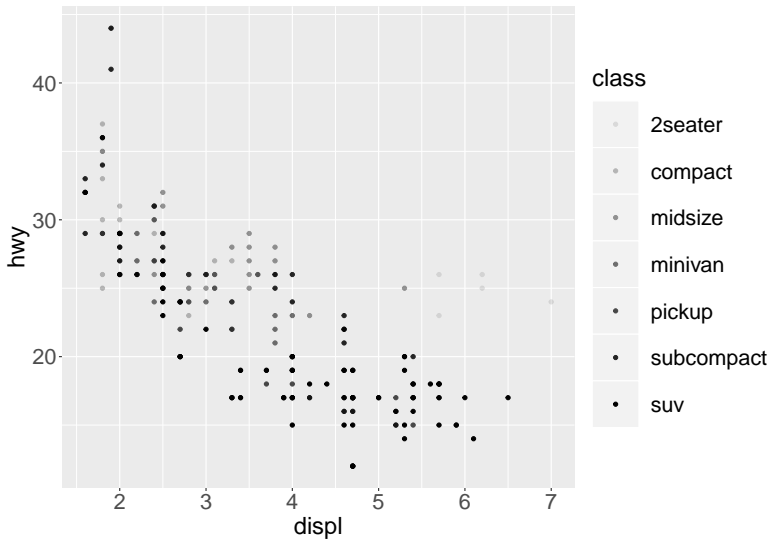
- ▶ This is because mapping an unordered variable (class) to ordered aesthetic (size) is not a good idea.
- ▶ Ordered aesthetics, like size and alpha, are more appropriate for continuous variables.

The alpha aesthetic

The alpha aesthetic changes the transparency of the points:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy, alpha = class))
```

The alpha aesthetic

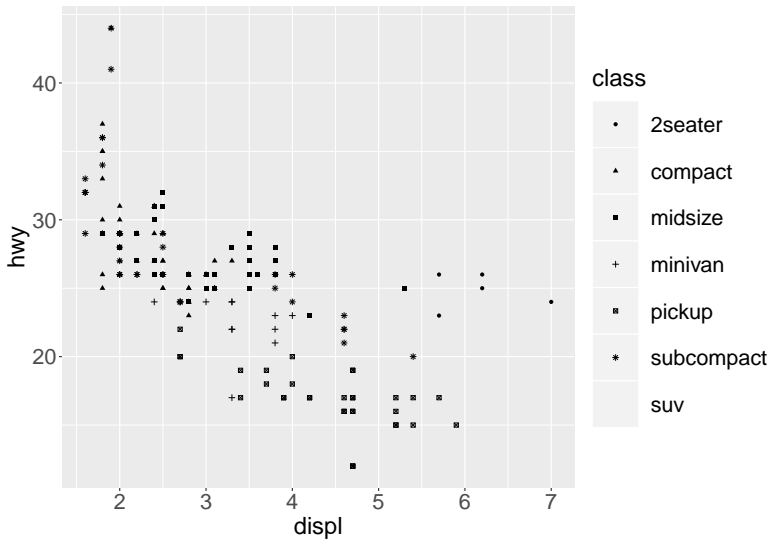


The shape aesthetic

The shape aesthetic changes the shape of the points:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy, shape = class))
```

The shape aesthetic



The shape aesthetic

- ▶ What happened to the SUVs?
- ▶ By default ggplot2 uses only up to six shapes at a time.
- ▶ This is because points become difficult to discriminate if we use too many shapes.
- ▶ We can, however, use more than six shapes if we set them “manually” with `scale_shape_manual()`.

The shape aesthetic

We can choose the following shapes of points by their number:





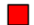







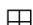







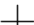




| | | | | |
|---|---|--|--|--|
|  0 |  4 |  10 |  15 |  22 |
|  1 |  6 |  11 |  16 |  21 |
|  2 |  7 |  12 |  17 |  24 |
|  5 |  8 |  13 |  18 |  23 |
|  3 |  9 |  14 |  19 |  20 |

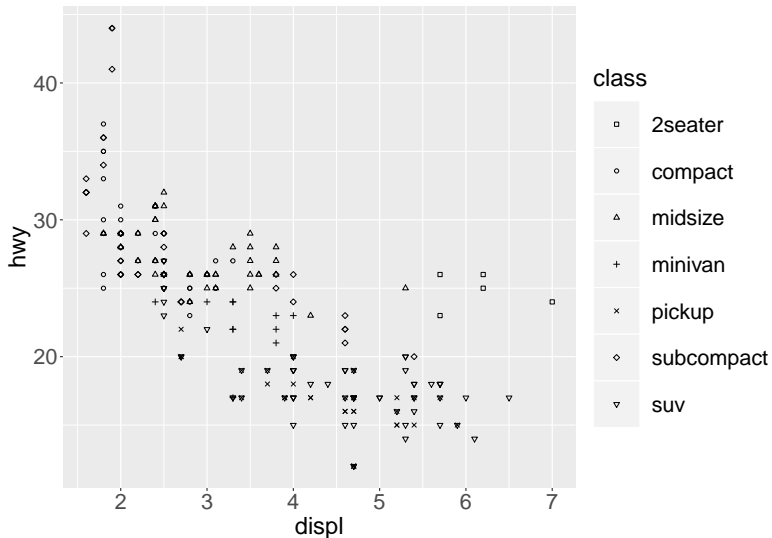
Figure 3: R Built in shapes

The shape aesthetic

Let's choose the first 7 shapes:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy, shape = class)) +  
  scale_shape_manual(values = 0:6)
```


The shape aesthetic

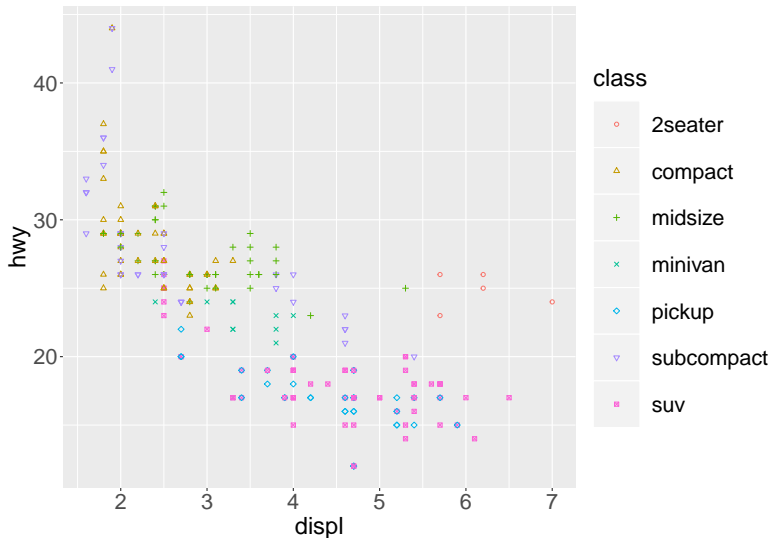


Mapping one variable to two aesthetics

- ▶ Mapping a single variable to multiple aesthetics is redundant, and should be avoided in most cases.
- ▶ In this case, however, since we have many categories, it can improve visual discrimination.
- ▶ Let's map class to color and shape:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy,  
                 shape = class, color = class)) +  
  scale_shape_manual(values = 1:7)
```

Mapping one variable to two aesthetics



Mapping continuous variables to aesthetics

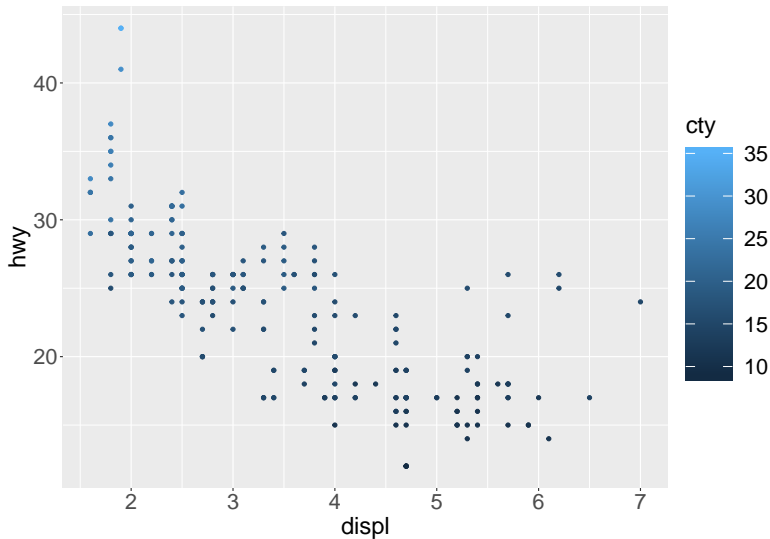
- ▶ So far we mapped `class`, a categorical variable, to the `color`, `size`, `shape` and `alpha` aesthetics of `geom_point()`.
- ▶ Now let's map continuous variables to these aesthetics and see how they behave differently for categorical versus continuous variables.
- ▶ The variable `cty` (city miles per gallon) is a continuous variable.

Mapping continuous variables to aesthetics

Let's map cty to the color aesthetic:

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy, colour = cty))
```

Mapping continuous variables to aesthetics

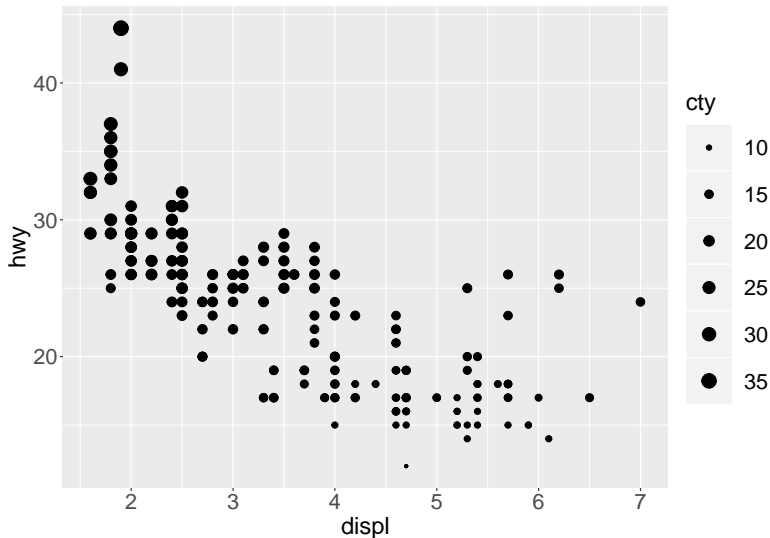


Mapping continuous variables to aesthetics

Now let's map `cty` to the size aesthetic:

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy, size = cty))
```

Mapping continuous variables to aesthetics




Mapping continuous variables to aesthetics

Now let's map `cty` to the `alpha` aesthetic:

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy, alpha = cty))
```

Mapping continuous variables to aesthetics



Mapping continuous variables to aesthetics

And finally let's map `cty` to the shape aesthetic:

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy, shape = cty))
```

```
#> Error: A continuous variable can not be mapped to  
#> shape
```

Aesthetic mappings: categorical vs continuous variables

Color and **alpha** aesthetics:

- ▶ Each level of a categorical variable is assigned to a different (discrete) color.
- ▶ The color of the points vary continuously from light to dark as a function of the values of a continuous variables.

Size aesthetic:

- ▶ Each level of a categorical variable is assigned to a different (discrete) size.
- ▶ The size of the points vary continuously as a function of the values of a continuous variables.

Aesthetic mappings: categorical vs continuous variables

Shape aesthetic:


- ▶ A different shape is assigned to each level of a categorical value.
- ▶ The shape aesthetic can not be mapped to continuous variables.

The linetype aesthetic

`geom_smooth()` draws a different type of line for each level of the variable that is mapped to the `linetype` aesthetic:

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv))
```

The linetype aesthetic



The linetype aesthetic

- ▶ The linetype aesthetic automatically adds a legend.
- ▶ The linetype aesthetic can not be mapped to continuous variables.

The group aesthetic

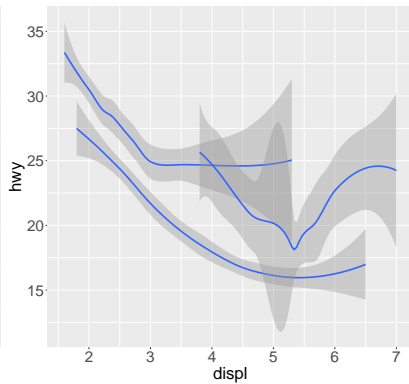
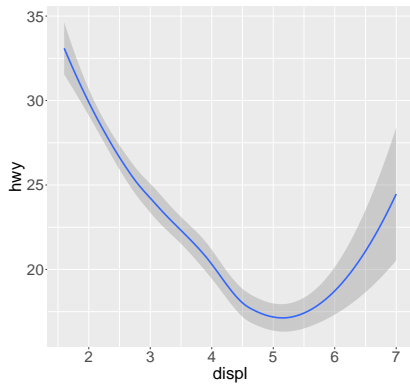
- ▶ The group aesthetic also draws a separate object for each level of a discrete variable.
- ▶ The group aesthetic does not add a legend or any other distinguishing features to the geoms.

The group aesthetic

```
p1 <- ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
p2 <- ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, group = drv))
```

The group aesthetic

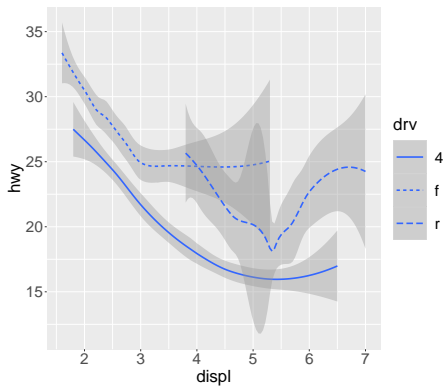
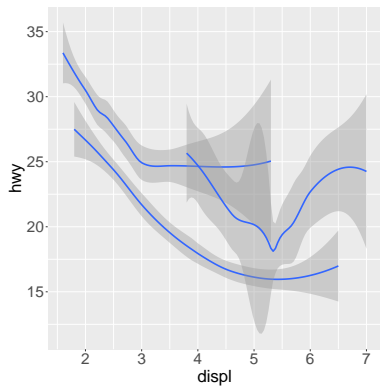


group vs linetype

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, group = drv))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv))
```

group vs linetype

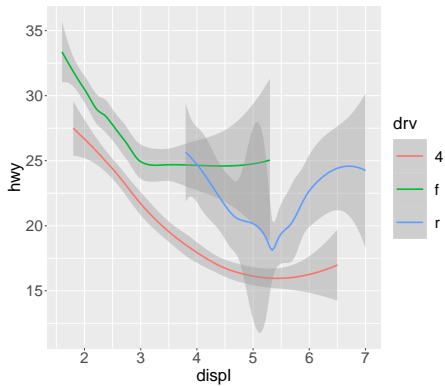
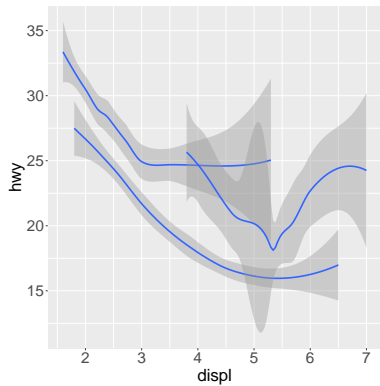


group vs color

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, group = drv))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, color = drv))
```

group vs color



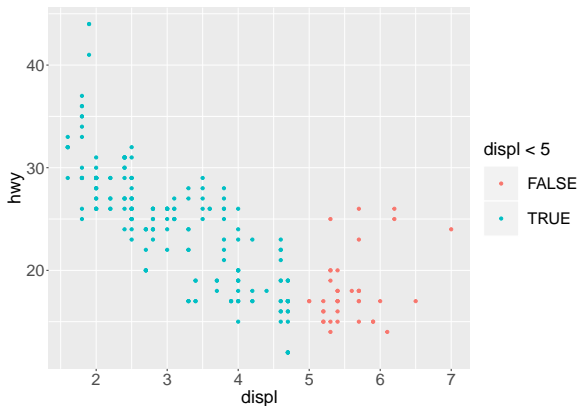
Aesthetics and logical conditions

Aesthetics can be mapped to logical expressions. For example, if you map an aesthetic to `displ < 5`:

- ▶ `ggplot()` creates a temporary variable with values equal to the result `displ < 5`.
- ▶ The result of `displ < 5` is a logical variable.
- ▶ `ggplot()` then maps aesthetics to the temporary variable.

Aesthetics and logical conditions

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy,  
                 colour = displ < 5))
```



Aesthetics and logical conditions

```
ggplot(mpg) +  
  geom_point(aes(x = displ, y = hwy,  
                 colour = class == "2seater"))
```



Changing default aesthetic properties of geoms

- ▶ To change a default aesthetic property, set the aesthetic by name as an argument of the geom function.
 - ▶ Names of colors should be indicated as character strings.
 - ▶ Size of points should be indicated in mm.
 - ▶ Shapes are identified by the numbers in figure 3.

Changing default aesthetic properties of geoms

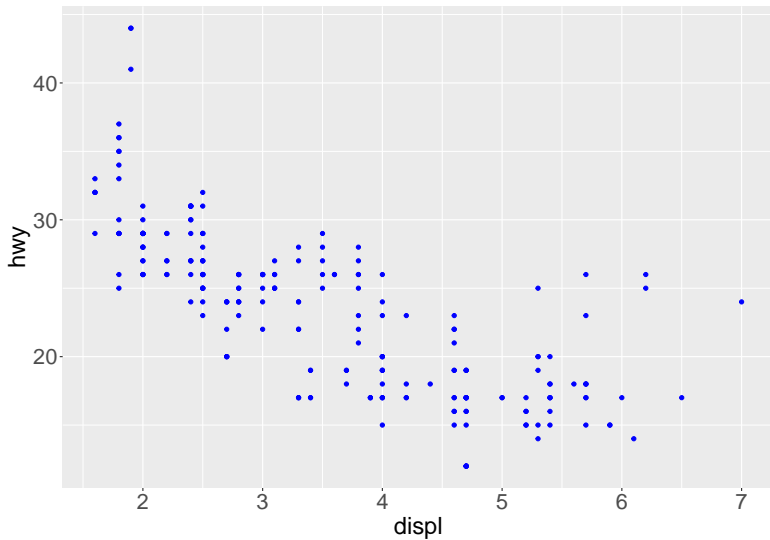
- ▶ There are some seeming duplicate shapes (e.g. 0, 15, and 22 are all squares).
- ▶ The difference comes from the interaction of the `colour` and `fill` aesthetics:
 - ▶ The hollow shapes (0–14) have a border determined by the `colour` aesthetic
 - ▶ The solid shapes (15–18) are filled with the `colour` aesthetic
 - ▶ The filled shapes (21–24) have a border set by the `colour` aesthetic and are filled with the `fill` aesthetic

Changing default aesthetic properties of geoms

For example, we can make all of the points blue:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "blue")
```

Changing default aesthetic properties of geoms




Changing default aesthetic properties of geoms

Now let's change the size of the points:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "blue",  
    size = 6)
```

Changing default aesthetic properties of geoms

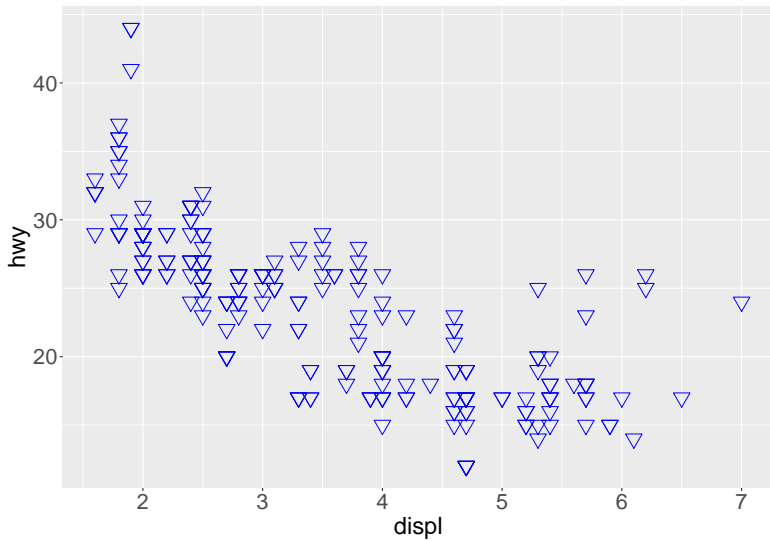


Changing default aesthetic properties of geoms

We can also change the default shape of the points:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "blue",  
    size = 5,  
    shape = 6)
```

Changing default aesthetic properties of geoms

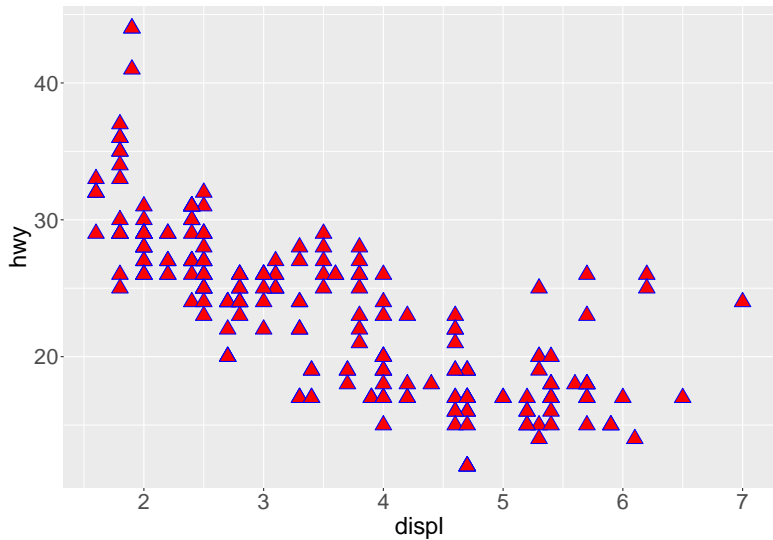


Changing default aesthetic properties of geoms

- ▶ Shape 6 is a hollow shape, it interacts only with the color aesthetic.
- ▶ If we want triangles filled with color, we must use shape 24, which interacts both with the color and fill aesthetics:

```
ggplot(data = mpg) +  
  geom_point(  
    mapping = aes(x = displ, y = hwy),  
    color = "blue",  
    size = 5,  
    shape = 24,  
    fill = "red"  
  )
```

Changing default aesthetic properties of geoms

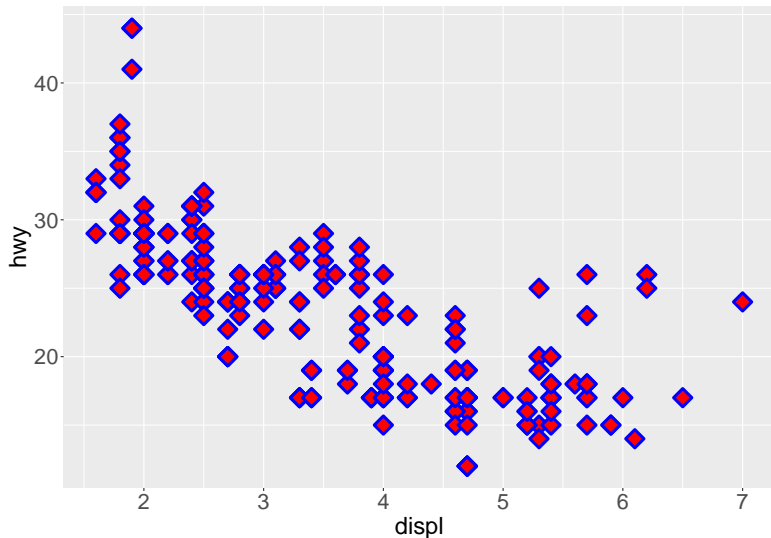


Changing default aesthetic properties of geoms

The stroke aesthetic changes the thickness of the border of the points:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "blue",  
    size = 5,  
    shape = 23,  
    fill = "red",  
    stroke = 2  
  )
```

Changing default aesthetic properties of geoms

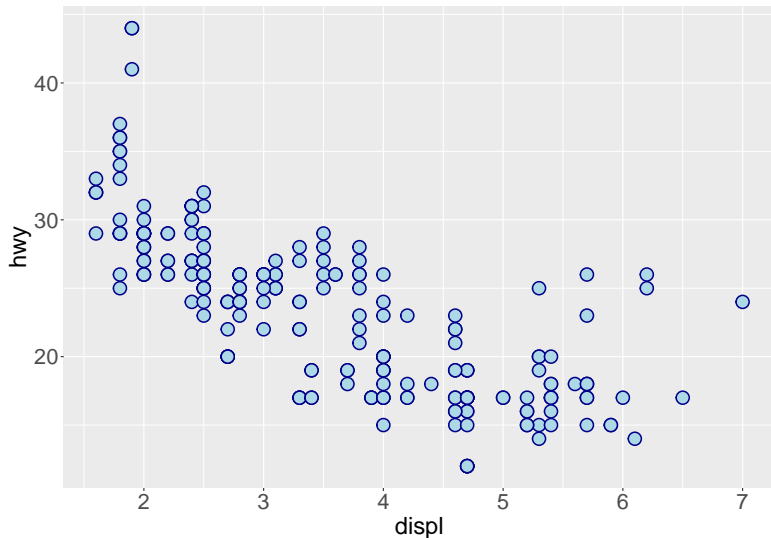


Changing default aesthetic properties of geoms

You can also use colors like dark blue and light blue:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "darkblue",  
    fill = "lightblue",  
    shape = 21,  
    size = 5,  
    stroke = 1  
  )
```

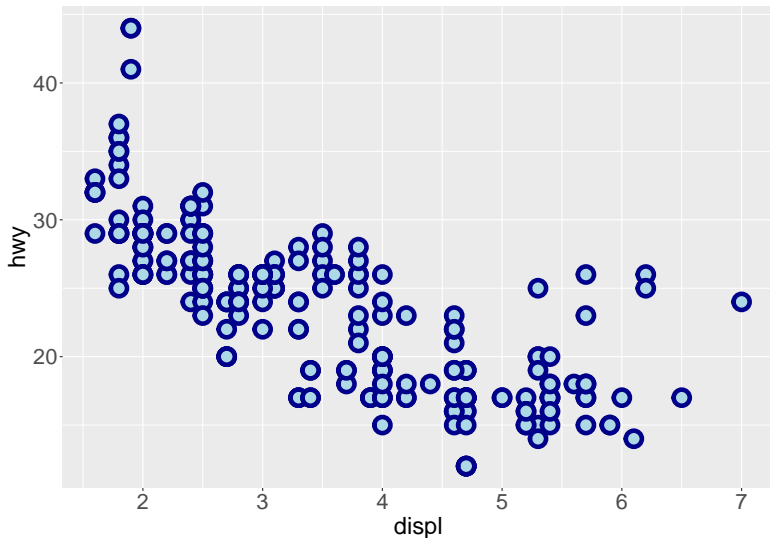
Changing default aesthetic properties of geoms



Changing default aesthetic properties of geoms

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
    color = "darkblue",  
    fill = "lightblue",  
    shape = 21,  
    size = 5,  
    stroke = 3  
  )
```

Changing default aesthetic properties of geoms



Changing default aesthetic properties of geoms

- For a list of available colors and their names see:
<http://sape.inf.usi.ch/quick-reference/ggplot2/colour>

Changing default aesthetic properties of geoms

- ▶ Alternatively, you can use RGB codes.
- ▶ The RGB model reproduces a broad array of colors by mixing together red, green, and blue in different proportions.
- ▶ You can specify colors like in HTML/CSS, using the hexadecimal values (00 to FF) for red, green, and blue, concatenated into a string, prefixed with a “#”.
- ▶ A pure red colour this is represented with “#FF0000”.

See:

- ▶ https://www.rapidtables.com/web/color/RGB_Color.html
- ▶ <https://htmlcolorcodes.com>

Changing default aesthetic properties of geoms

We can change the default line type of geom functions like `geom_smooth()`.



Figure 4: Built in line types

Changing default aesthetic properties of geoms

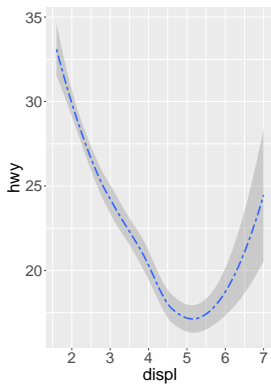
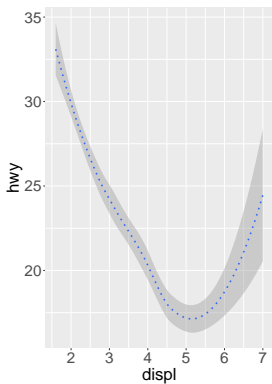
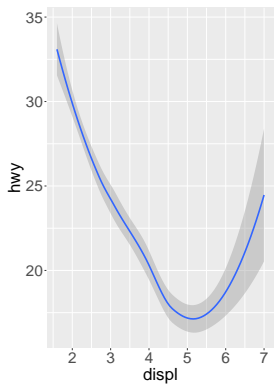
To change the default linetype, set it as an argument of the geom function:

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
    linetype = "dotted")
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
    linetype = "twodash")
```

Changing default aesthetic properties of geoms



Changing default aesthetic properties of geoms



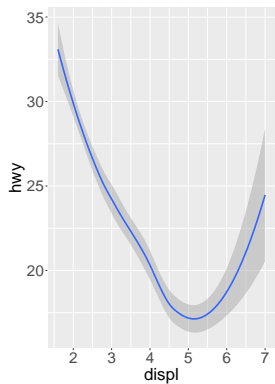
We can also change the thickness of the lines:

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy), size = 2)
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy), size = 5)
```


Changing default aesthetic properties of geoms



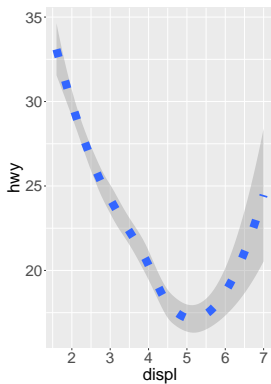
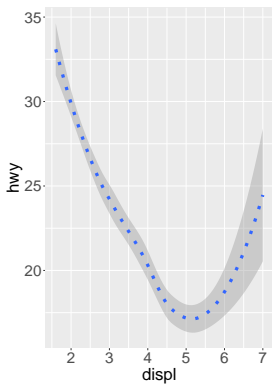
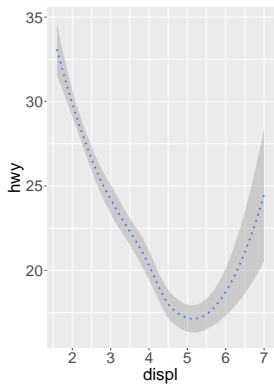
Changing default aesthetic properties of geoms

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
              linetype = "dotted")
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
              linetype = "dotted", size = 2)
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
              linetype = "dotted", size = 5)
```

Changing default aesthetic properties of geoms



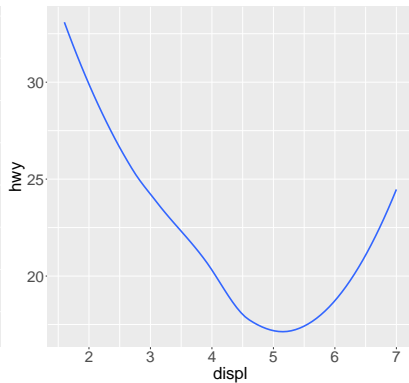
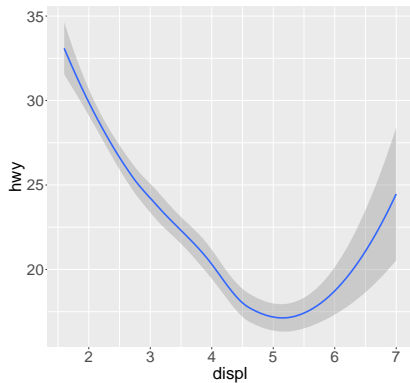
Changing default aesthetic properties of geoms

- ▶ By default, `geom_smooth()` displays confidence intervals around the smoothed line.
- ▶ Confidence intervals may be disabled by setting the `se` (standard error) aesthetic to `FALSE`:

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy),  
    se = FALSE)
```

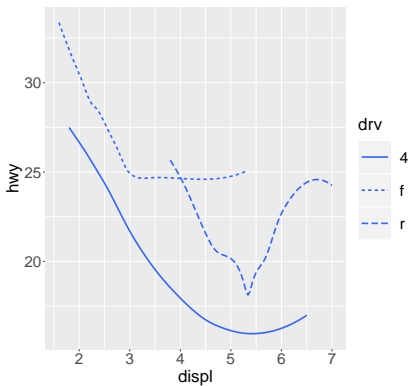
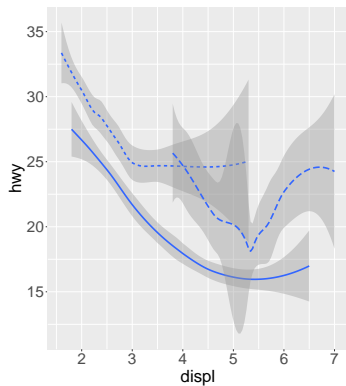
Changing default aesthetic properties of geoms



Changing default aesthetic properties of geoms

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv))  
  
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv),  
             se = FALSE)
```

Changing default aesthetic properties of geoms

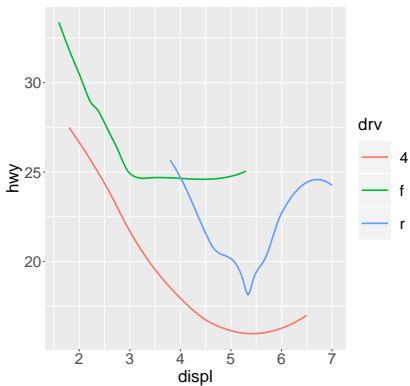
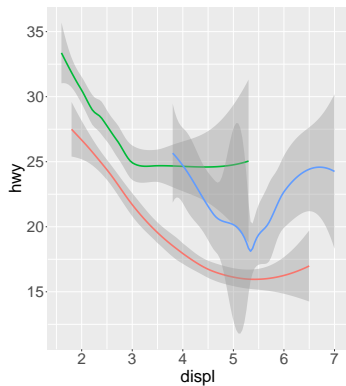


Changing default aesthetic properties of geoms

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, color = drv))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, color = drv),  
    se = FALSE)
```


Changing default aesthetic properties of geoms



Changing default aesthetic properties of geoms

- ▶ The default smoothing method of `geom_smooth` can be changed.
- ▶ The default method ("auto") depends on the size of the largest group.
- ▶ For less than 1000 observations the method is "loess" (a local polynomial regression).
- ▶ Admissible methods are: "auto", "lm", "glm", "gam" and "loess".

Changing default aesthetic properties of geoms

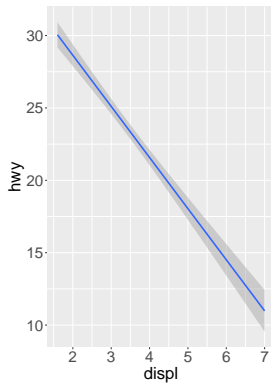
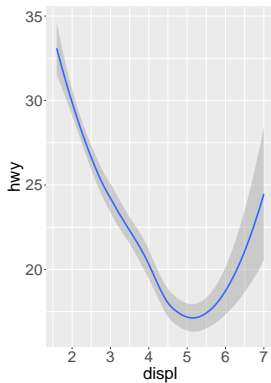
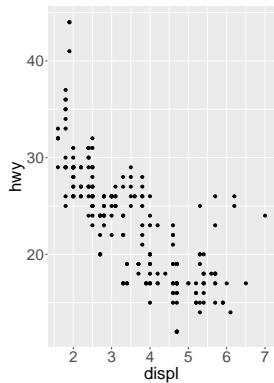
To add a linear regression line use `method = "lm"`:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy), method = "lm")
```

Changing default aesthetic properties of geoms




Changing default aesthetic properties of geoms

Legends can also be disabled:

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv),  
    se = FALSE)
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, linetype = drv),  
    se = FALSE, show.legend = FALSE)
```

Changing default aesthetic properties of geoms

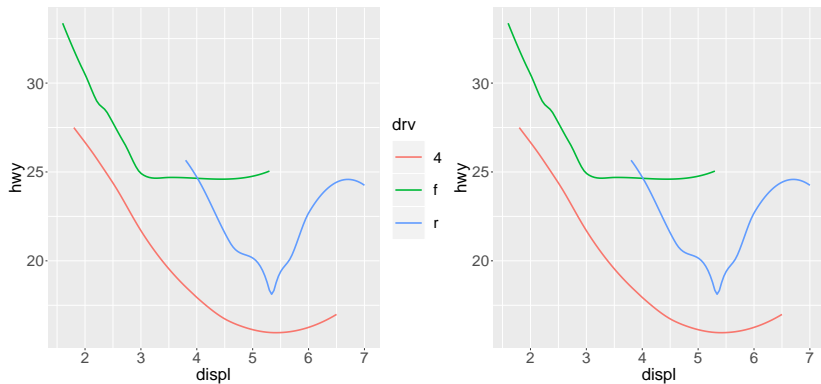


Changing default aesthetic properties of geoms

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, color = drv),  
    se = FALSE)
```

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, color = drv),  
    se = FALSE, show.legend = FALSE)
```

Changing default aesthetic properties of geoms



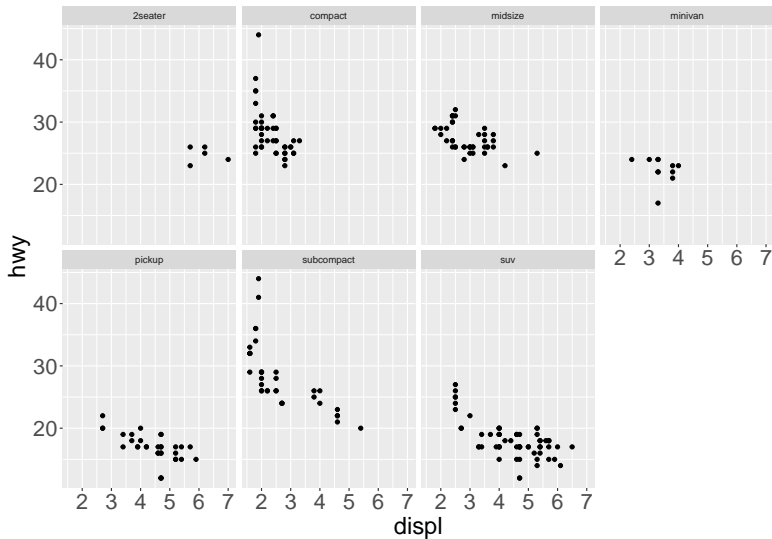
Facets

- ▶ One way to add additional variables to a plot is with **aesthetics**.
- ▶ Another way, particularly useful for categorical variables, is to split your plot into **facets**, subplots that each display one subset of the data.
- ▶ To facet your plot by a single variable, use `facet_wrap()`.
- ▶ The first argument of `facet_wrap()` should be "~" followed the name of a variable.

Facets

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  facet_wrap(~ class, nrow = 2)
```

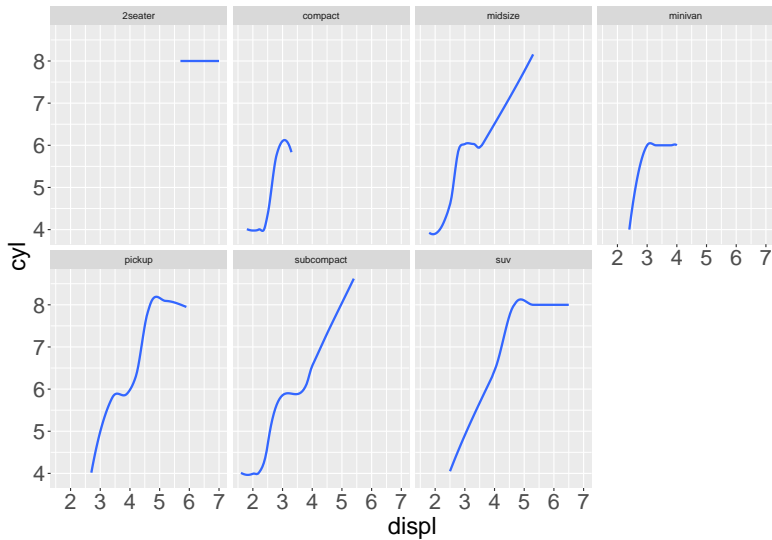
Facets



Facets

```
ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = cyl), se = FALSE) +  
  facet_wrap(~class, nrow = 2)
```

Facets

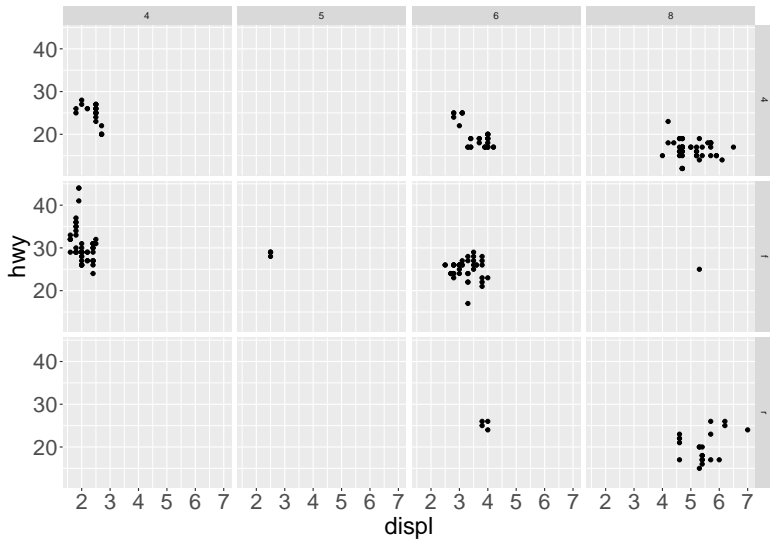


Facets

- ▶ To facet your plot on the combination of two variables, use `facet_grid()`.
- ▶ The first argument of `facet_grid()` should contain the two variable names separated by "~".

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  facet_grid(drv ~ cyl)
```

Facets

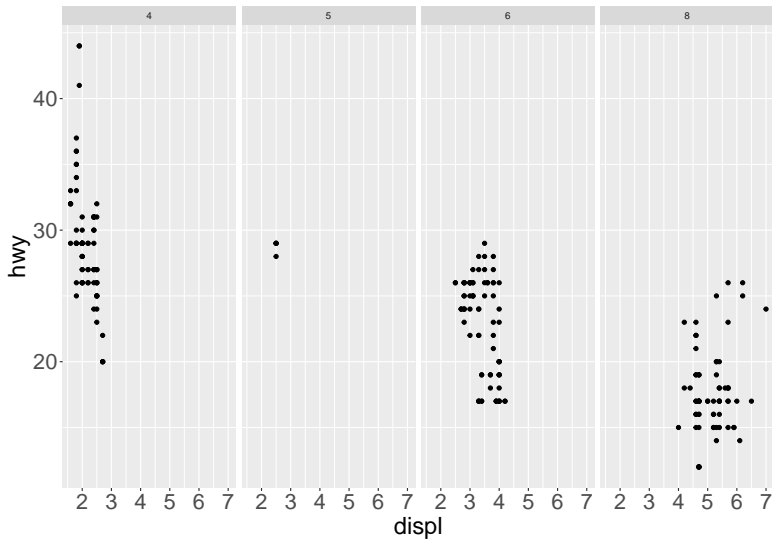


Facets

If you prefer to not facet in the rows (or columns), use a "." instead of a variable name for example:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  facet_grid(. ~ cyl)
```

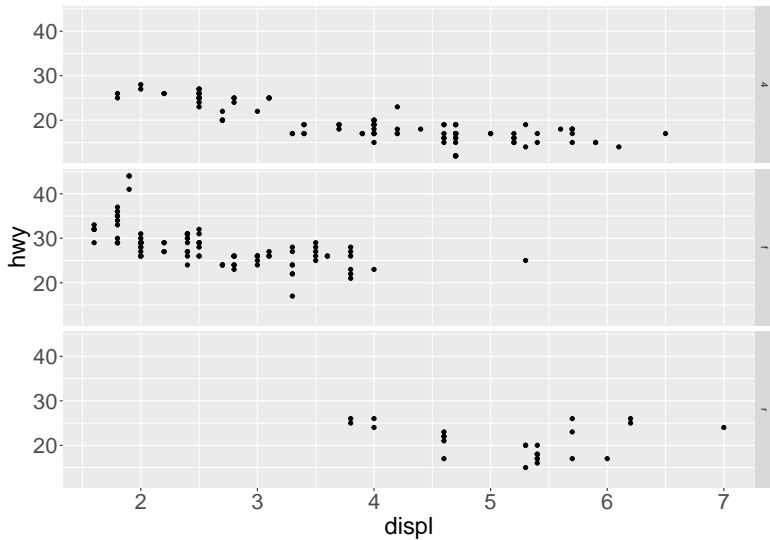

Facets



Facets

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  facet_grid(drv ~ .)
```

Facets

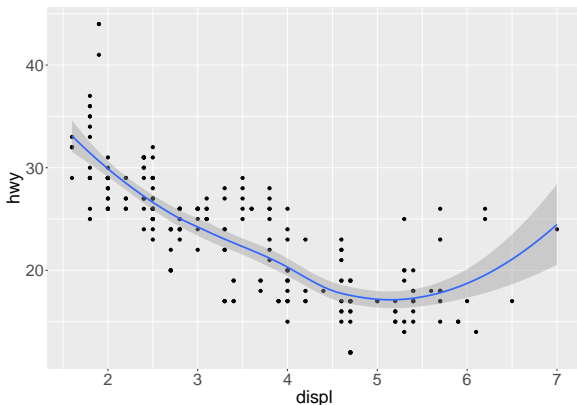


Multiple geoms

- ▶ We can have multiple geoms in the same graph.
- ▶ For example, we can overlap points and lines.
- ▶ This is achieved by adding multiple geom functions to `ggplot()`.

Multiple geoms

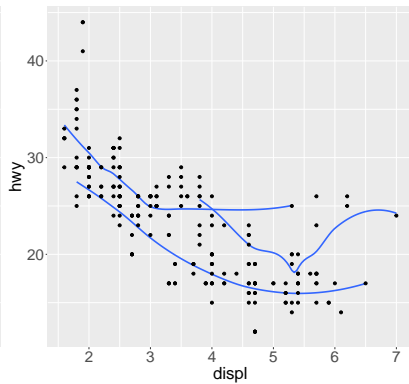
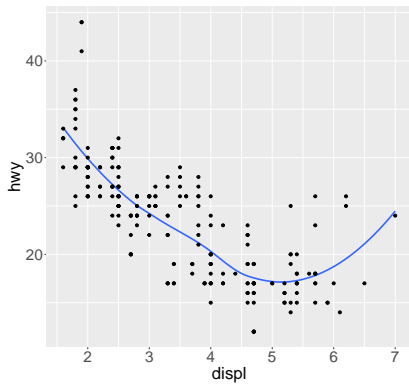
```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  geom_smooth(aes(x = displ, y = hwy))
```



Multiple geoms

```
p1 <- ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy), se = FALSE) +  
  geom_point(aes(x = displ, y = hwy))  
  
p2 <- ggplot(data = mpg) +  
  geom_smooth(aes(x = displ, y = hwy, group = drv),  
             se = FALSE) +  
  geom_point(aes(x = displ, y = hwy))  
  
ggarrange(p1, p2, nrow = 1)
```

Multiple geoms



Multiple geoms

- ▶ This, however, induces duplication in our code.
- ▶ We built the same aesthetic mapping twice in each graph.
- ▶ We can avoid this type of repetition by passing a set of mappings to `ggplot()`.
- ▶ `ggplot2` will treat these mappings as global mappings that apply to each geom in the graph.

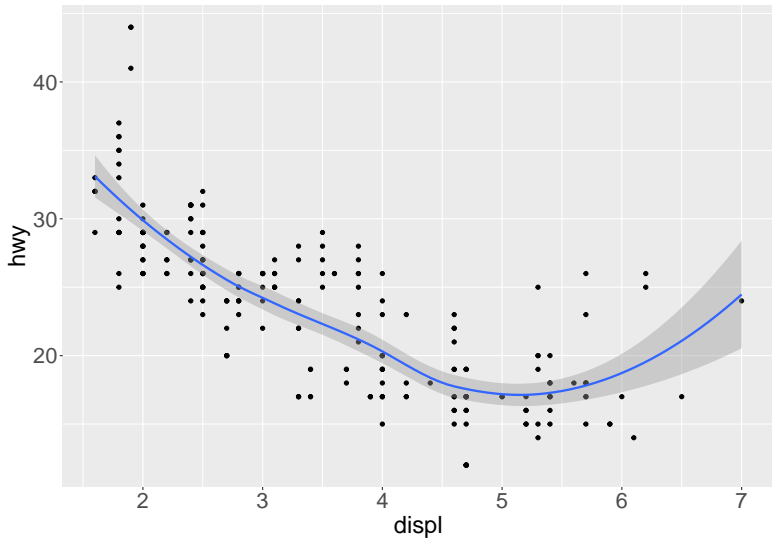
Multiple geoms

These two chunks of code result in the same plot:

```
ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy)) +  
  geom_smooth(aes(x = displ, y = hwy))
```

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()
```

Multiple geoms




Multiple geoms

- ▶ If you place mappings in a geom function, ggplot2 will treat them as local mappings for that layer.
- ▶ It will use these mappings to extend or overwrite the global mappings for that layer only.
- ▶ This makes it possible to display different aesthetics in different layers.

Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_smooth(aes(color = drv), se = FALSE) +  
  geom_point()
```

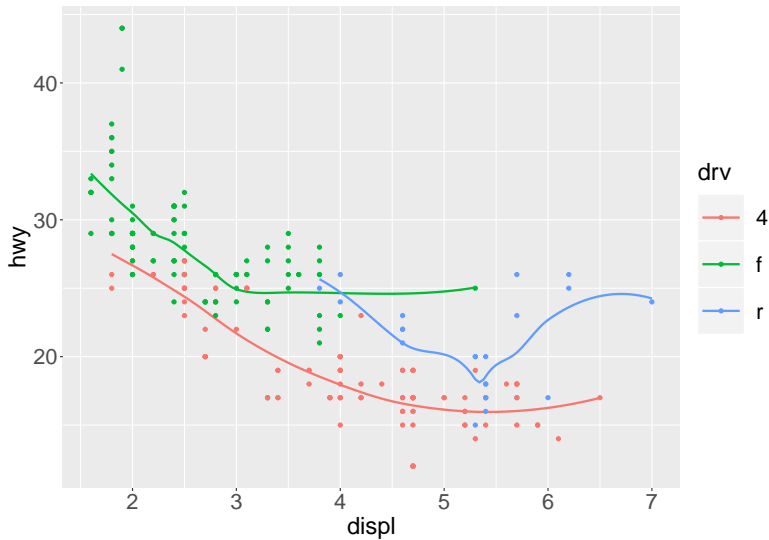
Multiple geoms



Multiple geoms

```
ggplot(data = mpg,  
       aes(x = displ, y = hwy, color = drv)) +  
  geom_smooth(se = FALSE) +  
  geom_point()
```

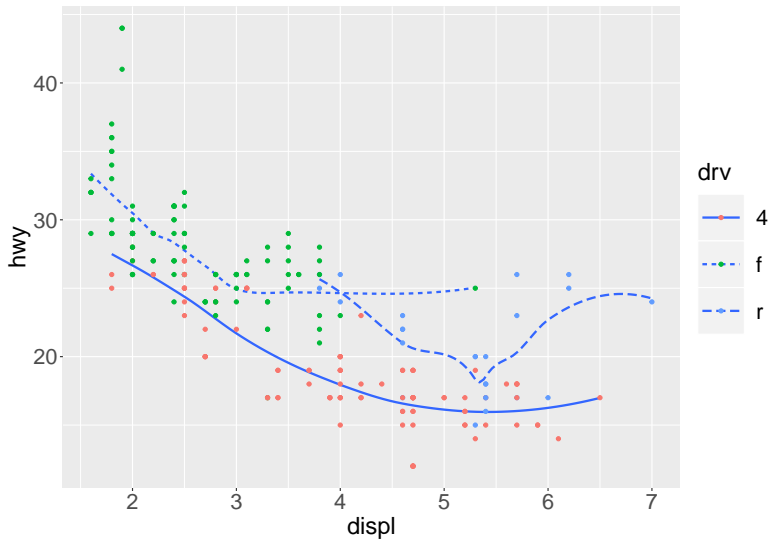
Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_smooth(aes(linetype = drv), se = FALSE) +  
  geom_point(aes(color = drv))
```



Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy,  
                        color = drv)) +  
  geom_smooth(aes(linetype = drv), se = FALSE) +  
  geom_point()
```

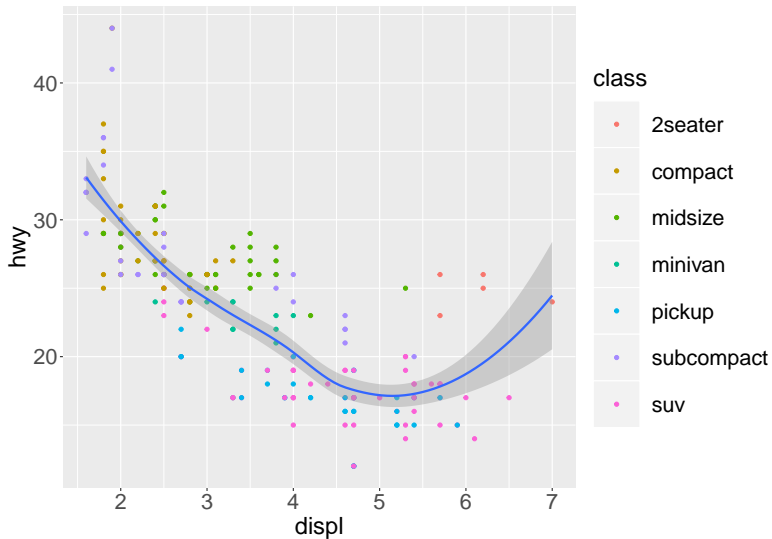
Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth()
```

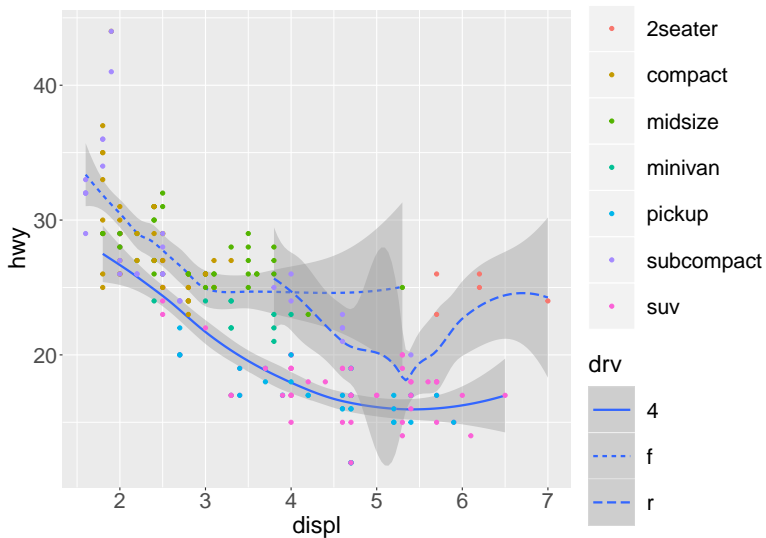
Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_smooth(aes(linetype = drv)) +  
  geom_point(aes(color = class))
```

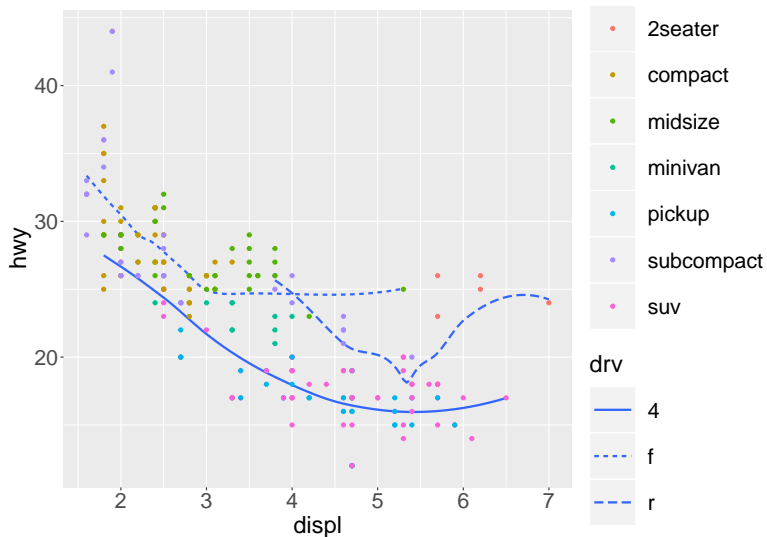
Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_smooth(aes(linetype = drv), se = FALSE) +  
  geom_point(aes(color = class))
```


Multiple geoms

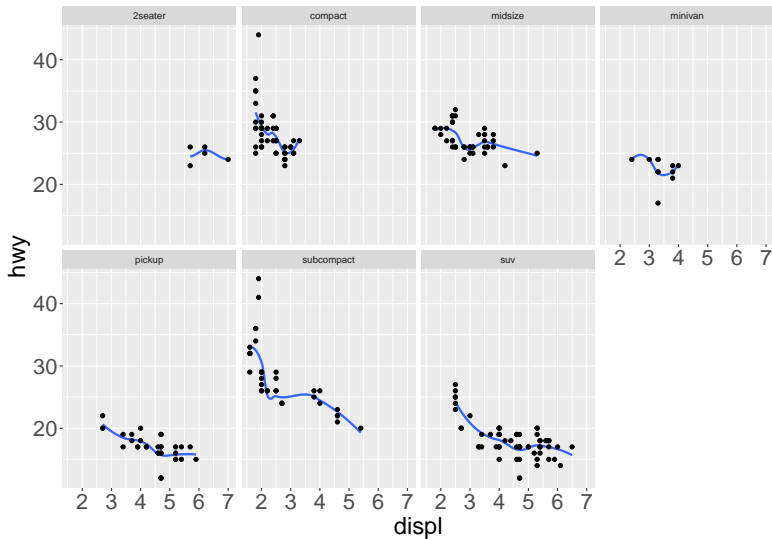


Multiple geoms

We can also overlap geoms in grids:

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_smooth(se = FALSE) +  
  geom_point() +  
  facet_wrap(~class, nrow = 2)
```

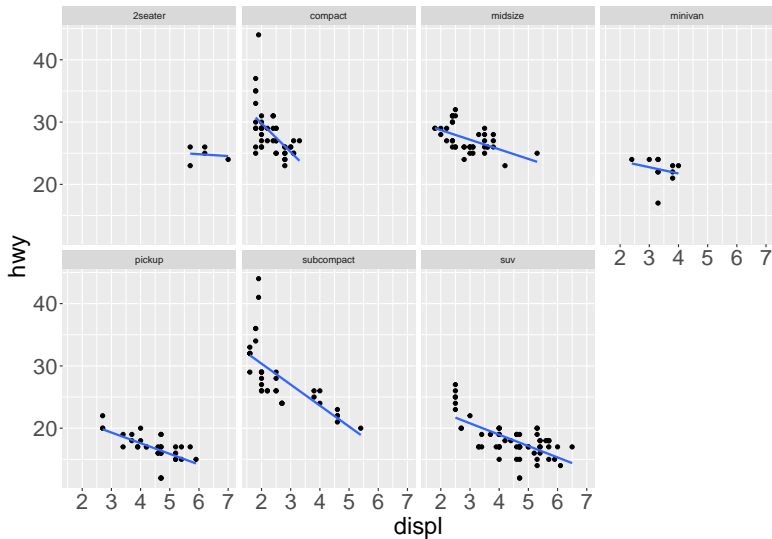
Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE) +  
  facet_wrap(~class, nrow = 2)
```

Multiple geoms



Multiple geoms

- ▶ You can use the same logic to specify different data for each layer.


Multiple geoms

In this example:

- ▶ The smooth line displays just a subset of the mpg dataset, the subcompact cars.
- ▶ The local data argument in `geom_smooth()` overrides the global data argument in `ggplot()` for that layer only.

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(  
    data = subset(mpg, class == "subcompact"),  
    se = FALSE  
  )
```


Multiple geoms



Multiple geoms

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(  
    data = subset(mpg, class == "subcompact"),  
    se = FALSE  
  ) +  
  geom_smooth(  
    data = subset(mpg, class == "2seater"),  
    se = FALSE  
  )
```

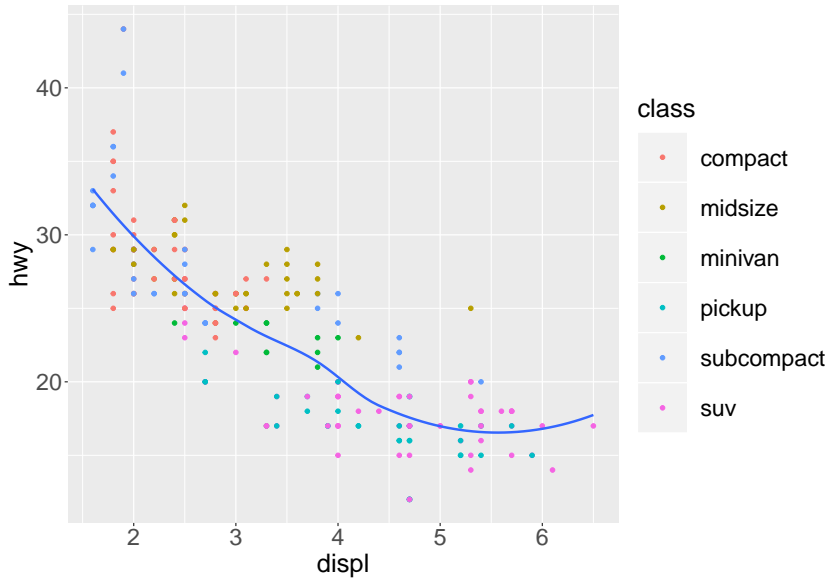
Multiple geoms



Multiple geoms

```
ggplot(data = subset(mpg, class != "2seater"),  
       aes(x = displ, y = hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(se = FALSE)
```

Multiple geoms



Multiple geoms

```
ggplot(data = subset(mpg, class != "2seater"),  
       aes(x = displ, y = hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(  
    data = subset(mpg,  
                  class %in% c("suv", "subcompact")),  
    aes(linetype = class), se = FALSE)
```

Multiple geoms




The diamonds dataset

- ▶ Diamonds is a dataset included in the ggplot2 package.
- ▶ Contains attributes of almost 54000 diamonds.
- ▶ The variables include price, carat, quality of the cut, color and clarity

Bar charts

Let's use a bar chart to display the number of diamonds for each quality of cut:

```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut))
```



Statistical transformations

- ▶ On the y-axis, bar charts displays counts.
- ▶ But counts are not a variable in diamonds dataset!
- ▶ Many graphs, like scatterplots, plot the raw values the dataset.
- ▶ Other graphs, like bar charts, calculate new values to plot.
- ▶ The algorithm used to calculate new values for a graph is called a **stat**, short for statistical transformation.

Statistical transformations

- ▶ Bar charts, histograms, and frequency polygons bin your data and then plot bin counts.
- ▶ Smoothers fit a model to your data and then plot predictions from the model.
- ▶ Boxplots compute summary statistics and then display them on specially formatted box.

Statistical transformations

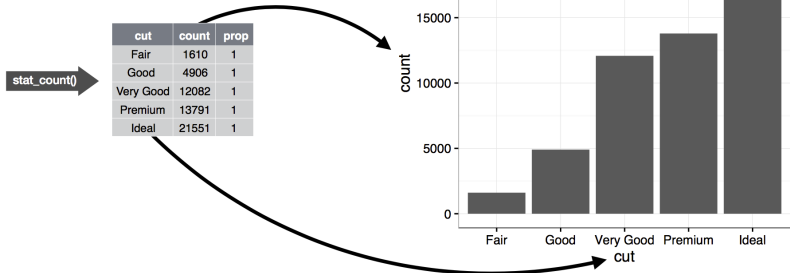
1. **geom_bar()** begins with the **diamonds** data set

[illegible]

Statistical transformations

2. **geom_bar()** transforms the data with the "count" stat, which returns a data set of cut values and counts.

3. **geom_bar()** uses the transformed data to build the plot. cut is mapped to the x axis, count is mapped to the y axis.




Statistical transformations

- ▶ You can learn which stat a geom function uses by inspecting the default value of the stat argument.
- ▶ For example, with `?geom_bar` we learn that the default stat of `geom_bar()` is `count`.
- ▶ This means that `geom_bar()` uses `stat_count()` as the default statistical transformation.
- ▶ You can generally use geoms and stats interchangeably.

Statistical transformations

For example, you can recreate the previous plot using `stat_count()` instead of `geom_bar()`:

```
ggplot(data = diamonds) +  
  stat_count(aes(x = cut))
```



Statistical transformations

This works because:

- ▶ Every geom has a default stat, and every stat has a default geom.
- ▶ The default stat of `geom_bar()` is `count`.
- ▶ The default geom of `stat_count()` is `bar`.

Overwriting the default stat

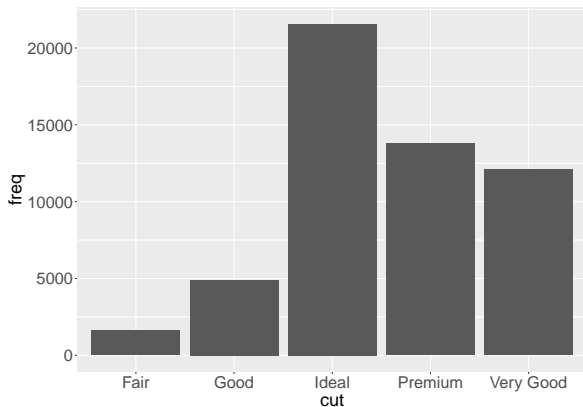
What if we want a bar chart that plot data as is?

- ▶ We may have a table with column heights.
- ▶ In that case, we need to change the default statistical transformation

```
tib <- tribble(  
  ~cut,      ~freq,  
  "Fair",    1610,  
  "Good",    4906,  
  "Very Good", 12082,  
  "Premium",  13791,  
  "Ideal",    21551  
)
```


Overriding the default stat

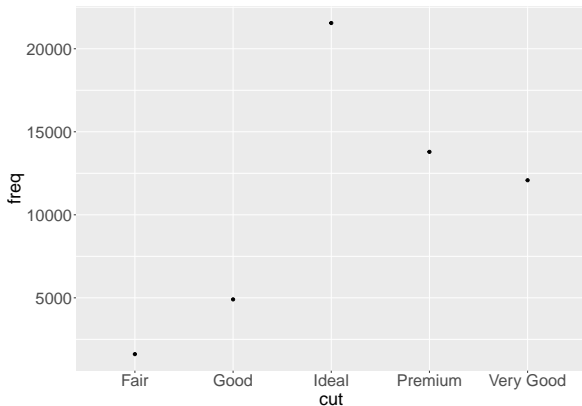
```
ggplot(data = tib) +  
  geom_bar(aes(x = cut, y = freq),  
    stat = "identity")
```



Overriding the default stat

The default stat of `stat_identity()` is point, not bar:

```
ggplot(data = tib) +  
  stat_identity(aes(x = cut, y = freq))
```

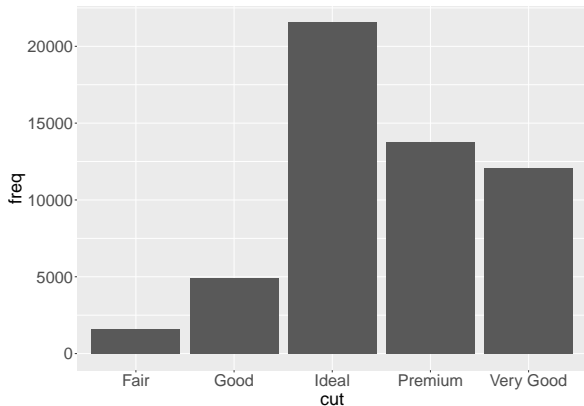


geom_col()

- ▶ To plot data as is, we can use `geom_col()`.
- ▶ The default stat of `geom_col()` is `stat_identity()`.
- ▶ `geom_col()` expects a column of y values with bar heights.

```
geom_col()
```

```
ggplot(data = tib) +  
  geom_col(aes(x = cut, y = freq))
```



Statistical transformations

- ▶ Stat functions calculate more variables than the ones that end up being displayed.
- ▶ To find the variables computed by a stat, look for the help section titled “computed variables”.
- ▶ From `?stat_count` we learn that the computed variables are counts and proportions.
- ▶ `ggplot_build()` let's us see every value that is calculated in the process of building a graph.

Statistical transformations

```
plt_1 <- ggplot(data = diamonds) +  
  geom_bar(aes(x = cut))
```

```
plt_1 <- ggplot_build(plt_1)  
plt_1$data[[1]][, 1:8]
```

| ## | y | count | prop | x | PANEL | group | ymin | ymax |
|------|-------|-------|------|---|-------|-------|------|-------|
| ## 1 | 1610 | 1610 | 1 | 1 | 1 | 1 | 0 | 1610 |
| ## 2 | 4906 | 4906 | 1 | 2 | 1 | 2 | 0 | 4906 |
| ## 3 | 12082 | 12082 | 1 | 3 | 1 | 3 | 0 | 12082 |
| ## 4 | 13791 | 13791 | 1 | 4 | 1 | 4 | 0 | 13791 |
| ## 5 | 21551 | 21551 | 1 | 5 | 1 | 5 | 0 | 21551 |

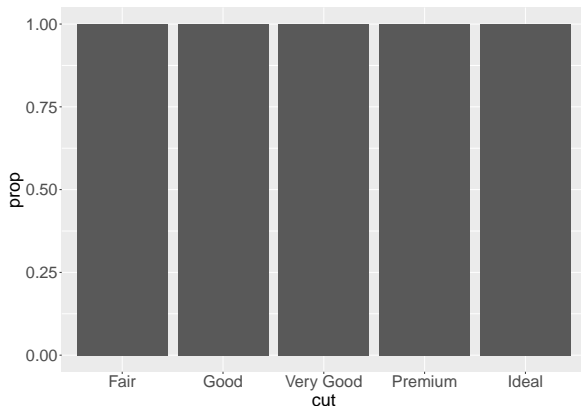
Mapping from transformed variables to aesthetics

- ▶ We can override the default mapping from transformed variables to aesthetics.
- ▶ For example, we might want to display a bar chart of proportions rather than counts.

Mapping from transformed variables to aesthetics

Let's map proportions, instead of counts, to the y axis:

```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, y = stat(prop)))
```



What went wrong?

Mapping from transformed variables to aesthetics

Lets see the computed values:

```
plt_2 <- ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, y = stat(prop)))  
  
plt_2 <- ggplot_build(plt_2)  
plt_2$data[[1]][, 1:8]
```

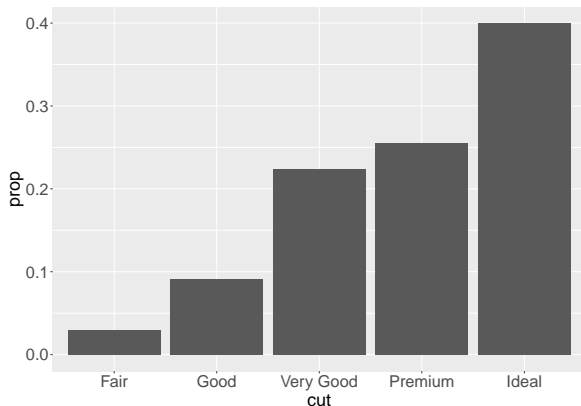
| ## | y | count | prop | x | PANEL | group | ymin | ymax |
|------|---|-------|------|---|-------|-------|------|------|
| ## 1 | 1 | 1610 | 1 | 1 | 1 | 1 | 0 | 1 |
| ## 2 | 1 | 4906 | 1 | 2 | 1 | 2 | 0 | 1 |
| ## 3 | 1 | 12082 | 1 | 3 | 1 | 3 | 0 | 1 |
| ## 4 | 1 | 13791 | 1 | 4 | 1 | 4 | 0 | 1 |
| ## 5 | 1 | 21551 | 1 | 5 | 1 | 5 | 0 | 1 |

Mapping from transformed variables to aesthetics

- ▶ The `prop` column is created as `count` divided by the sum of all of the counts that belong to the same group.
- ▶ By default, `ggplot2` created one group for each level of `x`, so:
 - ▶ Each proportion is calculated by dividing the count of each group by itself.
 - ▶ All the proportions are set to 1.
- ▶ To display proportions instead of counts we have to tell `ggplot2` that there is only one group so that it divides the counts by the total number of observations.

Mapping from transformed variables to aesthetics

```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, y = stat(prop), group = 1))
```



Mapping from transformed variables to aesthetics

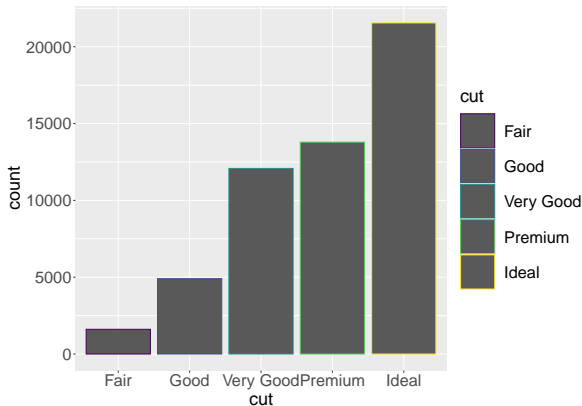
```
plt_3 <- ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, y = stat(prop), group = 1))  
plt_3 <- ggplot_build(plt_3)  
  
plt_3$data[[1]][, 1:5]
```

| ## | | y | count | prop | x | group |
|------|------------|-------|------------|------|---|-------|
| ## 1 | 0.02984798 | 1610 | 0.02984798 | 1 | 1 | |
| ## 2 | 0.09095291 | 4906 | 0.09095291 | 2 | 1 | |
| ## 3 | 0.22398962 | 12082 | 0.22398962 | 3 | 1 | |
| ## 4 | 0.25567297 | 13791 | 0.25567297 | 4 | 1 | |
| ## 5 | 0.39953652 | 21551 | 0.39953652 | 5 | 1 | |

Aesthetic of bar charts

You can map the color aesthetic to the grouping variable:


```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, color = cut))
```



Aesthetic of bar charts

We can also change the default fill color of `geom_bar()`:

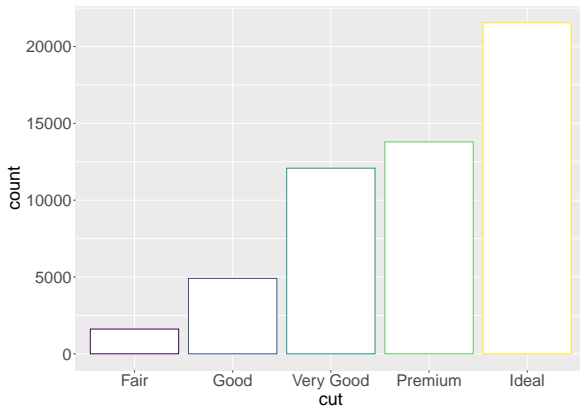
```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, color = cut), fill = "white")
```



Aesthetic of bar charts

And disable the legend:

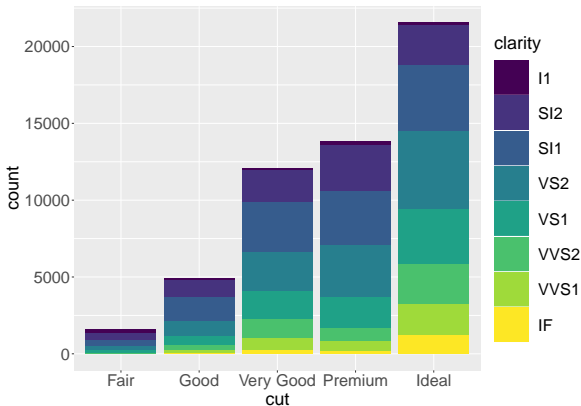
```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, color = cut),  
           fill = "white", show.legend = FALSE)
```



Stacked bar charts

The fill aesthetic can be mapped to variables other than x to add a dimension to the plot:

```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, fill = clarity))
```



Position adjustments

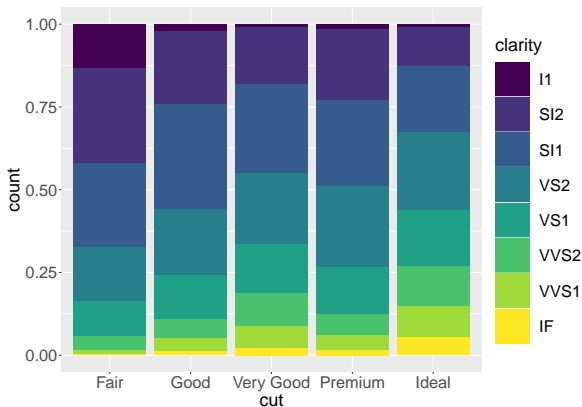
- ▶ Each colored rectangle represents a combination of cut and clarity.
- ▶ The stacking is performed automatically by the position adjustment specified in the `position` argument of the `geom` function.
- ▶ The default value is `stack`.

Position adjustments

- ▶ `position = "fill"` works like stacking but makes each set of stacked bars the same height.
- ▶ This makes it easier to compare proportions across groups.

Position adjustments


```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, fill = clarity),  
    position = "fill")
```



Position adjustments


Setting `position = "dodge"` places overlapping objects directly beside one another:

```
ggplot(data = diamonds) +  
  geom_bar(aes(x = cut, fill = clarity),  
    position = "dodge")
```



Position adjustments

Recall our first scatterplot:



Did you notice that the plot displays only 126 points, even though there are 234 observations in the dataset?

Position adjustments

- ▶ The values of `hwy` and `displ` are rounded, and many points overlap each other.
- ▶ This problem is known as overplotting.
- ▶ This makes it hard to see where the mass of the data is.
- ▶ Are the data points spread equally throughout the graph, or is there one special combination of `hwy` and `displ` that contains 109 values?



Position adjustments

- ▶ Setting `position = "jitter"` adds a small amount of random noise to each point, spreading the points.
- ▶ While it makes your graph less accurate at small scales, it makes your graph more revealing at large scales.
- ▶ `ggplot2` comes with a shorthand for `geom_point(position = "jitter")`:
 - ▶ `geom_jitter()`.

Position adjustments

```
p1 <- ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy),  
             position = "jitter")  
  
p2 <- ggplot(data = mpg) +  
  geom_point(aes(x = displ, y = hwy))  
  
ggarrange(p1, p2, nrow = 1)
```


Position adjustments



Position adjustments



```
p1 <- ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point(position = "jitter") +  
  geom_smooth()
```

```
p2 <- ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()
```

```
ggarrange(p1, p2, nrow = 1)
```

Position adjustments

The confidence interval of the smoothed lines can also help:



Position adjustments

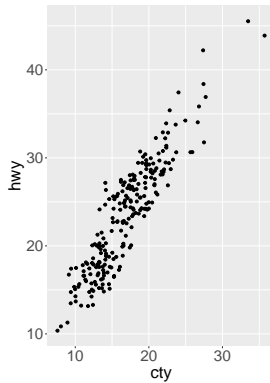
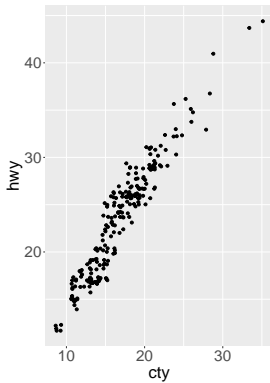
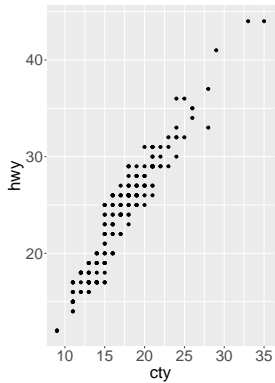
There are two optional arguments to jitter:

- ▶ Width controls the amount of vertical displacement.
- ▶ Height controls the amount of horizontal displacement.

Position adjustments


```
p1 <- ggplot(data = mpg, aes(x = cty, y = hwy)) +  
  geom_point()  
  
p2 <- ggplot(data = mpg, aes(x = cty, y = hwy)) +  
  geom_jitter()  
  
p3 <- ggplot(data = mpg, aes(x = cty, y = hwy)) +  
  geom_jitter(height = 2, width = 2)  
  
ggarrange(p1, p2, p3, nrow = 1)
```

Position adjustments



Boxplots


```
ggplot(data = mpg, aes(x = class, y = hwy)) +  
  geom_boxplot()
```



Coordinate systems

`coord_flip()` flips the coordinate system:


```
ggplot(data = mpg, aes(x = class, y = hwy)) +  
  geom_boxplot() +  
  coord_flip()
```



Coordinate systems


A pie chart is a stacked bar chart with polar coordinates:

```
ggplot(mpg, aes(x = "", fill = drv)) +  
  geom_bar()
```




Coordinate systems

```
ggplot(mpg, aes(x = "", fill = drv)) +  
  geom_bar() +  
  coord_polar("y")
```




Coordinate systems

```
ggplot(mpg, aes(x = "", fill = class)) +  
  geom_bar() +  
  coord_polar("y")
```



Coordinate systems

```
ggplot(mpg, aes(x = "", fill = drv)) +  
  geom_bar() +  
  coord_polar("y")
```



The ggplot template

```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(  
    mapping = aes(<MAPPINGS>),  
    stat = <STAT>,  
    position = <POSITION>  
  ) +  
  <COORDINATE_FUNCTION> +  
  <FACET_FUNCTION>
```

Bibliography

Wickham, Hadley, and Garrett Grolemund. 2016. "R for Data Science." *O'Reilly*.

Wickham, Hadley. 2010. "A Layered Grammar of Graphics." *Journal of Computational and Graphical Statistics* 19 (1). Taylor & Francis: 3–28.